

# Computer Systems Security

## CS 628A

Pramod Subramanyan  
Indian Institute of Technology Kanpur

# **MODULE 1: CONTROL HIJACKING**

# Acknowledgements

- Suman Jana
- Sandeep Shukla
- Dan Boneh (Stanford University)
- John C. Mitchell (Stanford University)
- Nicolai Zeldovich (MIT)
- Jungmin Park (Virginia Tech)
- Patrick Schaumont (Virginia Tech)
- Web Resources



# Control Hijacking

---

## Basic Control Hijacking Attacks

# Example 1

```
#include <stdio.h>
int main()
{
    volatile int number = 0;
    char buffer[8];

    printf("number=%d\n", number);
    gets(buffer);
    printf("number=%d\n", number);
    return 0;
}
```

```
$ python -c 'print "a"*8' | ./stack0
$ python -c 'print "a"*9' | ./stack0
```

## Example 2

```
int main(int argc, char* argv[])
{
    volatile int number = 0;
    char buffer[16];
    if (argc != 2) return 0;
    printf("number=%d\n", number);
    sprintf(buffer, "hello %s", argv[1]);
    printf("buffer=%s\n", buffer);
    printf("number=%d\n", number);
    return 0;
}
```

```
* $ ./stack1 `python -c 'print "a"*16'`
* $ ./stack1 `python -c 'print "a"*17'`
```

# Example 3

```
#include <stdio.h>
#include <string.h>
int checkpwd(const char* pwd)
{
    int good = 0;
    char buffer[8];
    strcpy(buffer, pwd);
    if(strcmp(buffer, "blahblah") == 0) good = 1;
    if(strcmp(buffer, "wahhwahh") == 0) good = 1;
    return good;
}
int main(int argc, char* argv[])
{
    if (argc != 2) return 1;
    if (checkpwd(argv[1])) {
        printf("=====\n");
        printf("|| access granted ||\n");
        printf("=====\n");
    } else { printf("denied!\n"); }
    return 0;
}
```

How do we exploit this?

# Example 4

```
#include <stdio.h>
#include <stdlib.h>

void win()
{
    printf("you win!\n");
}

int main(int argc, char* argv[])
{
    char buffer[16];

    printf("win addr=%p\n", win);
    gets(buffer);
    printf("buffer=%s\n", buffer);

    return 0;
}
```

How do we exploit this?



# Example 5

```
#include <stdio.h>

int main(int argc, char* argv[])
{
    volatile int number = 0;
    char *buffer = new char[16];
    if (argc != 2) return 0;

    printf("number=%d\n", number);
    sprintf(buffer, "hello %s", argv[1]);
    printf("buffer=%s\n", buffer);
    printf("number=%d\n", number);

    delete [] buffer;
    return 0;
}
```

How do we exploit this?

# Example 6

See Heap1.cpp

How do we exploit this?

# Resources


- The Hacking Book
  - Linked to in Moodle
  - Read the chapter on exploitation
  - These examples are based on code in the book
- Examples posted on course website
  - Run them on your computer
  - Play around with them in gdb: look at the addresses, step through them instr by instr
  - Do this now, and enjoy stress-free CTFs

# Control hijacking attacks

- Attacker's goal:
  - Take over target machine (e.g. web server)
    - Execute arbitrary code on target by hijacking application control flow
- Examples.
  - Buffer overflow attacks
  - Integer overflow attacks
  - Format string vulnerabilities

# Example 1: buffer overflows

- Extremely common bug in C/C++ programs.
  - First major exploit: 1988 Internet Worm. fingerd.



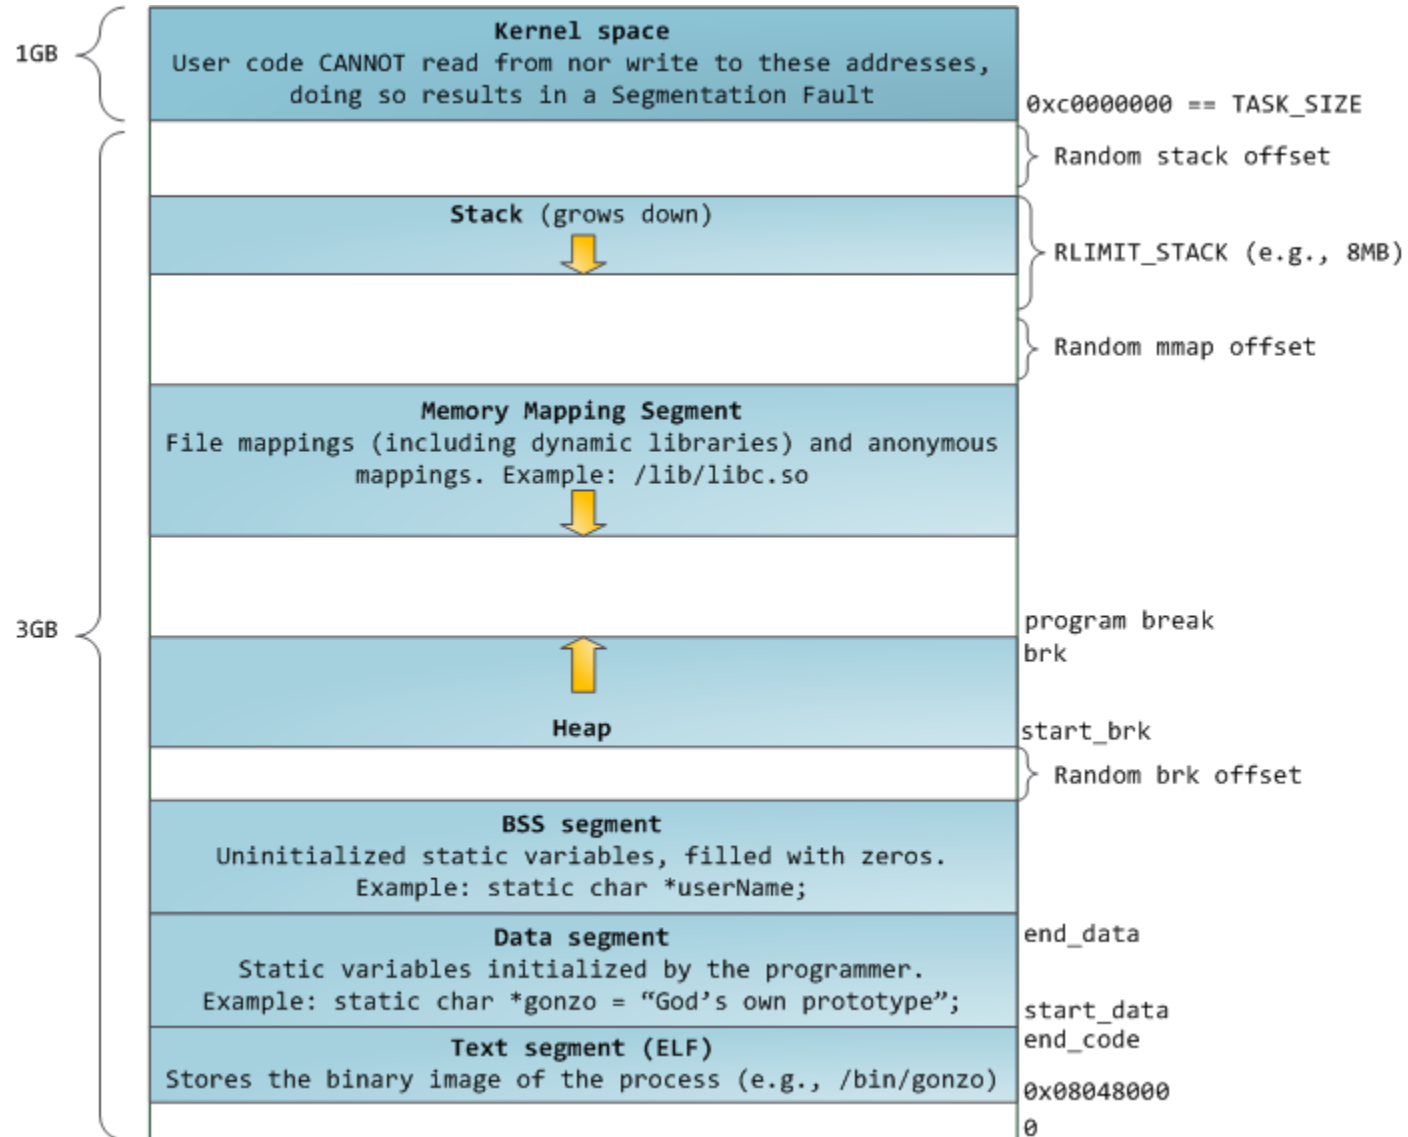
≈20% of all vuln.

Source: NVD/CVE

# What is needed

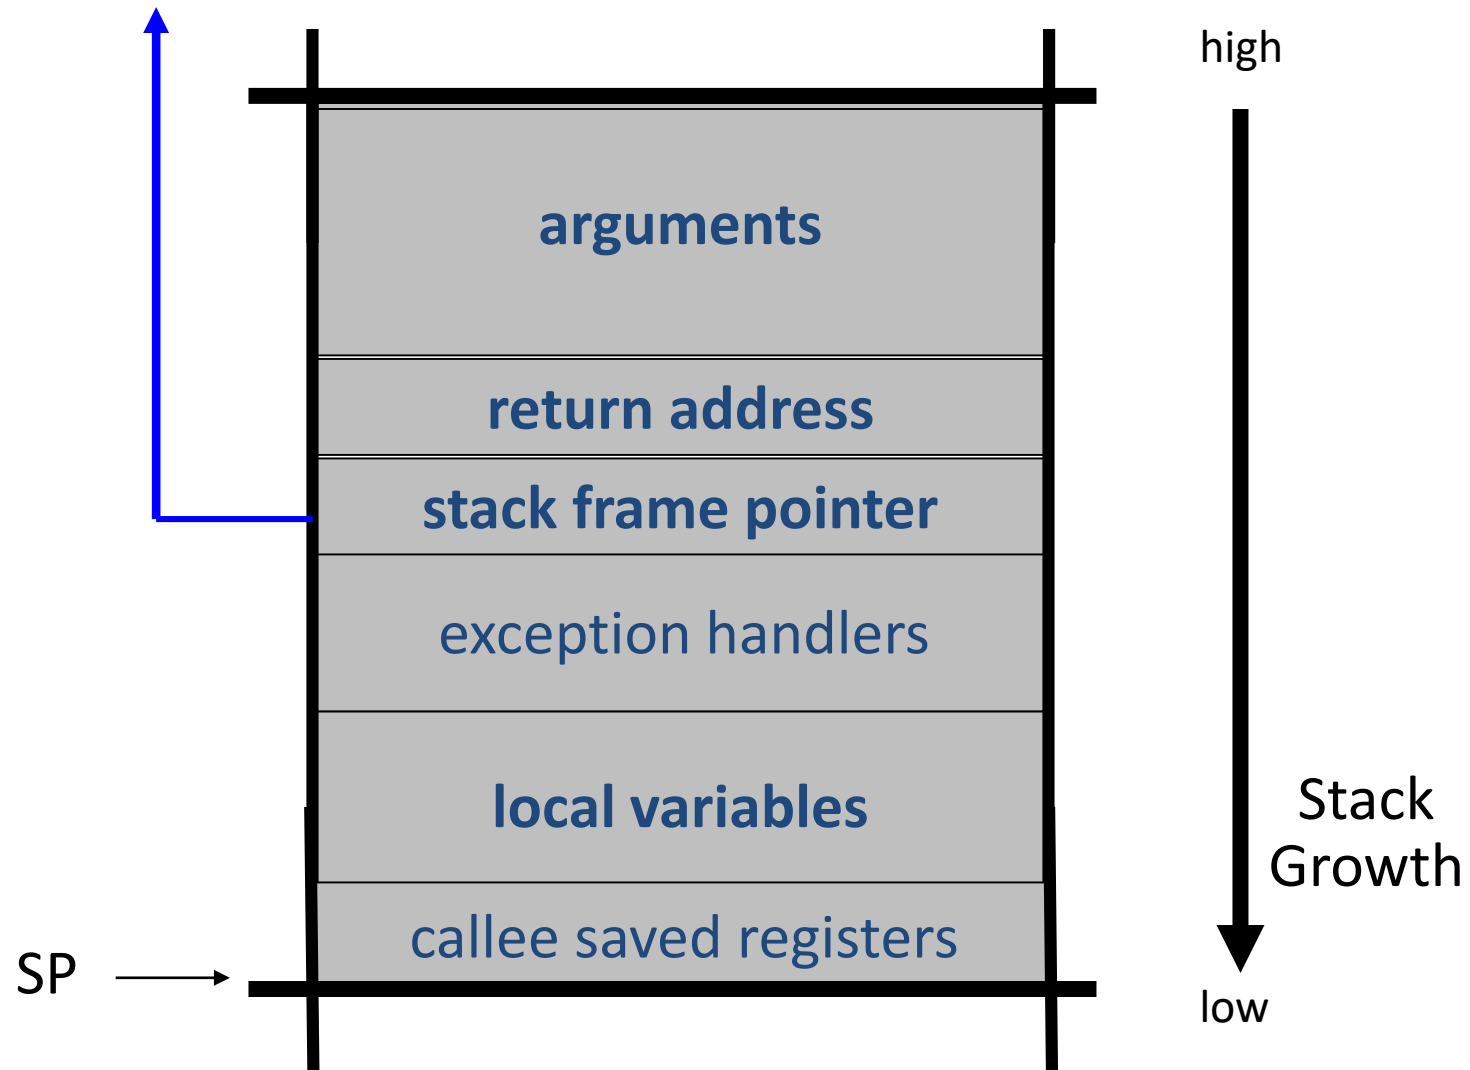
- Understanding C functions, the stack, and the heap.
  - Know how system calls are made
  - The `exec()` system call
- 
- Attacker needs to know which CPU and OS used on the target machine:
    - Our examples are for x86 running Linux or Windows
    - Details vary slightly between CPUs and OSs:
      - Little endian vs. big endian (x86 vs. Motorola)
      - Stack Frame structure (Unix vs. Windows)

# Linux Process Memory Layout



# Stack Frame

<http://post.queensu.ca/~trd/377/tut5/stack.html>



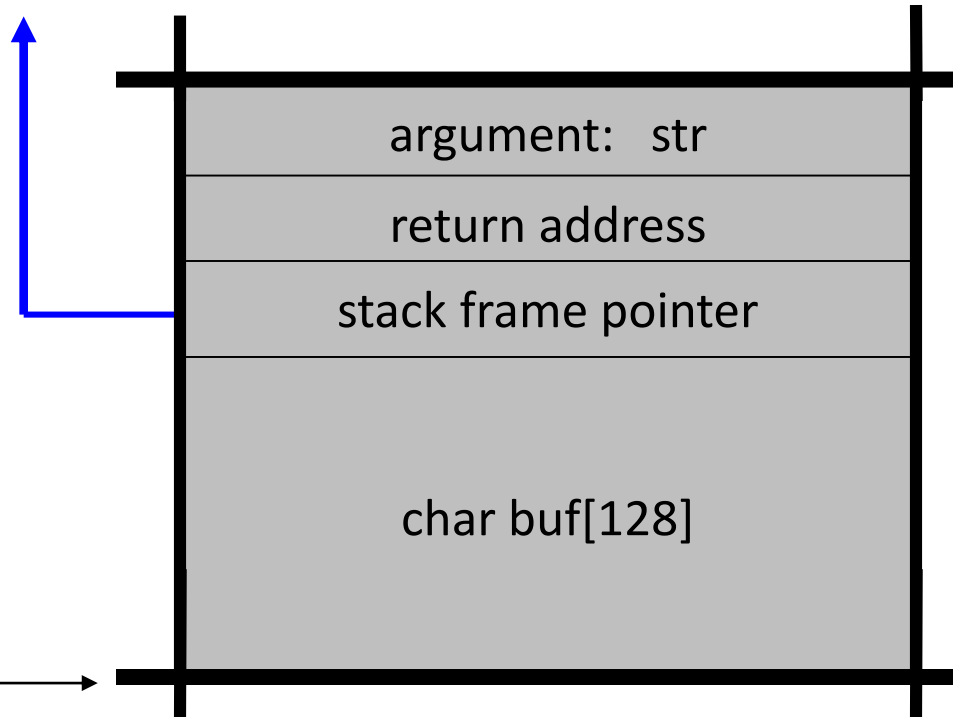


# What are buffer overflows?

Suppose a web server contains a function:

When func() is called stack looks like:

```
void func(char *str) {  
    char buf[128];  
  
    strcpy(buf, str);  
    do-something(buf);  
}
```




# What are buffer overflows?

What if `*str` is 136 bytes long?

After `strcpy`:

```
void func(char *str) {  
    char buf[128];  
  
    strcpy(buf, str);  
    do-something(buf);  
}
```



Problem:  
no length checking in `strcpy()`

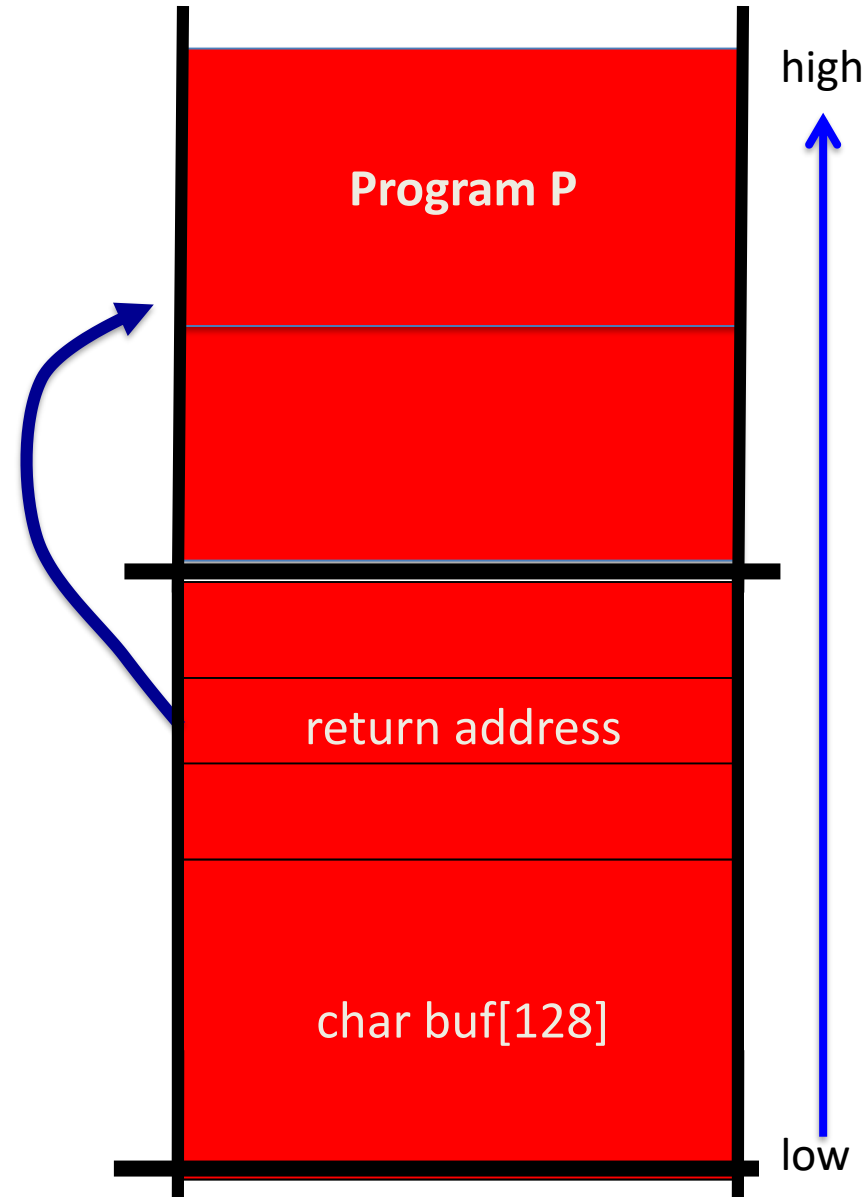
# Basic stack exploit

Suppose `*str` is such that  
after `strcpy` stack looks like:

Program P: `exec("/bin/sh")`

When `func()` exits, the user gets shell !

Note: attack code P runs *in stack*.



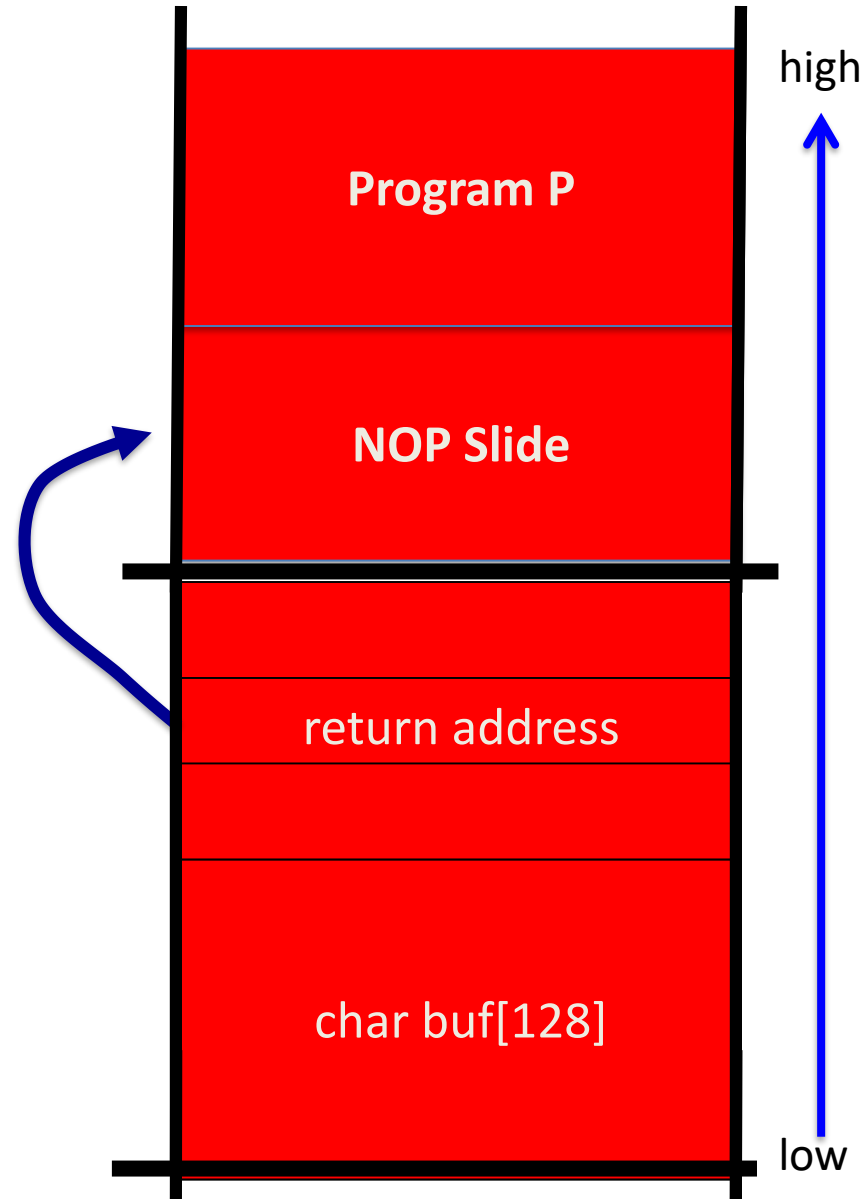
# The NOP slide

Problem: how does attacker determine ret-address?

Solution: NOP slide

- Guess approximate stack state when `func()` is called
- Insert many NOPs before program P:

`nop , xor eax,eax , inc ax`



# Details and examples

- Some complications:
  - Program P should not contain the '\0' character.
  - Overflow should not crash program before func() exists.
- (in)Famous remote stack smashing overflows:
  - (2007) Overflow in Windows animated cursors (ANI).  
LoadAniIcon() <https://www.sans.org/reading-room/whitepapers/threats/ani-vulnerability-history-repeats-1926>
  - (2005) Overflow in Symantec Virus Detection  
`test.GetPrivateProfileString "file", [long string]`

# Many unsafe libc functions

`strcpy` (char \*dest, const char \*src)

`strcat` (char \*dest, const char \*src)

`gets` (char \*s)

`scanf` ( const char \*format, ... )      and many more.

- “Safe” libc versions `strncpy()`, `strncat()` are misleading
  - e.g. `strncpy()` may leave string unterminated.
- Windows C run time (CRT):
  - `strcpy_s (*dest, DestSize, *src)`: ensures proper termination

# Buffer overflow opportunities


- Exception handlers: (Windows SEH attacks)
  - Overwrite the address of an exception handler in stack frame.
- Function pointers: (e.g. PHP 4.0.2, MS MediaPlayer Bitmaps)




- Overflowing buf will override function pointer.
- Longjmp buffers: longjmp(pos) (e.g. Perl 5.003)
- Overflowing buf next to pos overrides value of pos.

# Corrupting method pointers

- Compiler generated function pointers (e.g. C++ code)



- After overflow of **buf**:





# Finding buffer overflows

- To find overflow:
  - Run web server on local machine
  - Issue malformed requests (ending with “\$\$\$\$\$” )
    - Many automated tools exist (called fuzzers)
  - If web server crashes,  
search core dump for “\$\$\$\$\$” to find overflow location
- Construct exploit (not easy given latest defenses)



# Control Hijacking

---

## More Control Hijacking Attacks

# More Hijacking Opportunities

- **Integer overflows:** (e.g. MS DirectX MIDI Lib)
- **Double free:** double free space on heap
  - Can cause memory mgr to write data to specific location
  - Examples: CVS server
- **Use after free:** using memory after it is freed
- **Format string vulnerabilities**

# Integer Overflows

(see Phrack 60)

Problem: what happens when int exceeds max value?

**int m; (32 bits)**

**short s; (16 bits)**

**char c; (8 bits)**

$$c = 0x80 + 0x80 = 128 + 128$$

$$\Rightarrow c = 0$$

$$s = 0xff80 + 0x80$$

$$\Rightarrow s = 0$$

$$m = 0xffffffff80 + 0x80$$

$$\Rightarrow m = 0$$

Can this be exploited?

# An example


```
void func( char *buf1, *buf2,  unsigned int len1, len2) {  
    char temp[256];  
    if (len1 + len2 > 256) {return -1}           // length check  
    memcpy(temp, buf1, len1);                     // cat buffers  
    memcpy(temp+len1, buf2, len2);  
    do-something(temp);                           // do stuff  
}
```

What if **len1 = 0x80, len2 = 0xffffffff80** ?

⇒ len1+len2 = 0

Second memcpy() will overflow heap !!

# Integer overflow exploit stats



Source: NVD/CVE

# Format string bugs

# Format String Example 1

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int A = 5, B = 7, count_one, count_two;

    // Example of a %n format string
    printf("The number of bytes written up to this
point X%n is being stored in count_one, and the
number of bytes up to here X%n is being stored in
count_two.\n", &count_one, &count_two);

    printf("count_one: %d\n", count_one);
    printf("count_two: %d\n", count_two);

    // Stack Example
    printf("A is %d and is at %08x. B is %x.\n", A,
&A, B);

    exit(0);
}
```

\$ ./a.out

The number of bytes written up to this point X is being stored in count\_one, and the number of bytes up to here X is being storied in count\_two.

count\_one: 46

count\_two: 113

A is 5 and is at bffff7f4. B is 7.



# Format String Example 2

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int A = 5, B = 7, count_one, count_two;

    // Example of a %n format string
    printf("The number of bytes written up to this point X%n is
being stored in count_one, and the number of bytes up to here X%n
is being stored in count_two.\n", &count_one, &count_two);

    printf("count_one: %d\n", count_one);
    printf("count_two: %d\n", count_two);

    // Stack Example
    printf("A is %d and is at %08x. B is %x.\n", A, &A);

    exit(0);
}
```

\$ ./a.out

The number of bytes written up to this point X is being stored in count\_one, and the number of bytes up to here X is being storied in count\_two.

count\_one: 46

count\_two: 113

A is 5 and is at bffff7f4. B is b7fd6ff4

# Format String Example 3

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[]) {
    char text[1024];
    static int test_val = -72;

    if(argc < 2) {
        printf("Usage: %s <text to print>\n", argv[0]);
        exit(0);
    }
    strcpy(text, argv[1]);

    printf("The right way to print user-controlled input:\n");
    printf("%s", text);

    printf("\nThe wrong way to print user-controlled input:\n");
    printf(text);

    printf("\n");

    // Debug output
    printf("[*] test_val @ 0x%08x = %d 0x%08x\n", &test_val, test_val,
test_val);

    exit(0);
}
```

\$ ./fmt\_vuln testing%x

\$ ./fmt\_vuln \$(perl -e 'print "%08x."x40')

# Format string problem

```
int func(char *user) {  
    fprintf( stderr, user);  
}
```

Problem: what if \*user = "%s%s%s%s%s%s%s" ??

- Most likely program will crash: DoS.
- If not, program will print memory contents. Privacy?

Correct form:     `fprintf( stdout, "%s", user);`

# Vulnerable functions

Any function using a format string.

Printing:

printf, fprintf, sprintf, ...

vprintf, vfprintf, vsprintf, ...

Logging:

syslog, err, warn

# Exploit

- Dumping arbitrary memory:
  - Walk up stack until desired pointer is found.
  - `printf( "%08x.%08x.%08x.%08x|%s|")`
- Writing to arbitrary memory:
  - `printf( "hello %n", &temp)` -- writes '6' into temp.
  - `printf( "%08x.%08x.%08x.%08x.%n")`



# Control Hijacking

---

## Platform Defenses

# Preventing hijacking attacks

1. Fix bugs:
  - Audit software
    - Automated tools: Coverity, Prefast/Prefix.
  - Rewrite software in a type safe language (Java, ML)
    - Difficult for existing (legacy) code ...
2. Concede overflow, but prevent code execution
3. Add runtime code to detect overflows exploits
  - Halt process when overflow exploit detected
  - StackGuard, LibSafe, ...

# Marking memory as non-execute (W^X)


Prevent attack code execution by marking stack and heap as **non-executable** – **DEP** – **Data Execution Prevention**

- NX-bit on AMD Athlon 64, XD-bit on Intel P4 Prescott
  - NX bit in every Page Table Entry (PTE)
- Deployment:
  - Linux (via PaX project); OpenBSD
  - Windows: since XP SP2 (DEP)
    - Visual Studio: **/NXCompat[:NO]**
- Limitations:
  - Some apps need executable heap (e.g. JITs).
  - Does not defend against **Return Oriented Programming** exploits



# Attack: Return Oriented Programming (ROP)

- Control hijacking without executing code



# Response: randomization

- **ASLR**: (Address Space Layout Randomization)
  - Map shared libraries to rand location in process memory
    - ⇒ Attacker cannot jump directly to exec function
  - **Deployment**: (/DynamicBase)
    - **Windows 7**: 8 bits of randomness for DLLs
      - aligned to 64K page in a 16MB region ⇒ 256 choices
    - **Windows 8**: 24 bits of randomness on 64-bit processors
- **Other randomization methods**:
  - Sys-call randomization: randomize sys-call id's
  - Instruction Set Randomization (ISR)

# ASLR Example

Booting twice loads libraries into different locations:

ntlanman.dll	0x6D7F0000	Microsoft® Lan Manager
ntmarta.dll	0x75370000	Windows NT MARTA provider
ntshrui.dll	0x6F2C0000	Shell extensions for sharing
ole32.dll	0x76160000	Microsoft OLE for Windows

ntlanman.dll	0x6DA90000	Microsoft® Lan Manager
ntmarta.dll	0x75660000	Windows NT MARTA provider
ntshrui.dll	0x6D9D0000	Shell extensions for sharing
ole32.dll	0x763C0000	Microsoft OLE for Windows

Note: everything in process memory must be randomized  
**stack, heap, shared libs, base image**

- Win 8 **Force ASLR**: ensures all loaded modules use ASLR



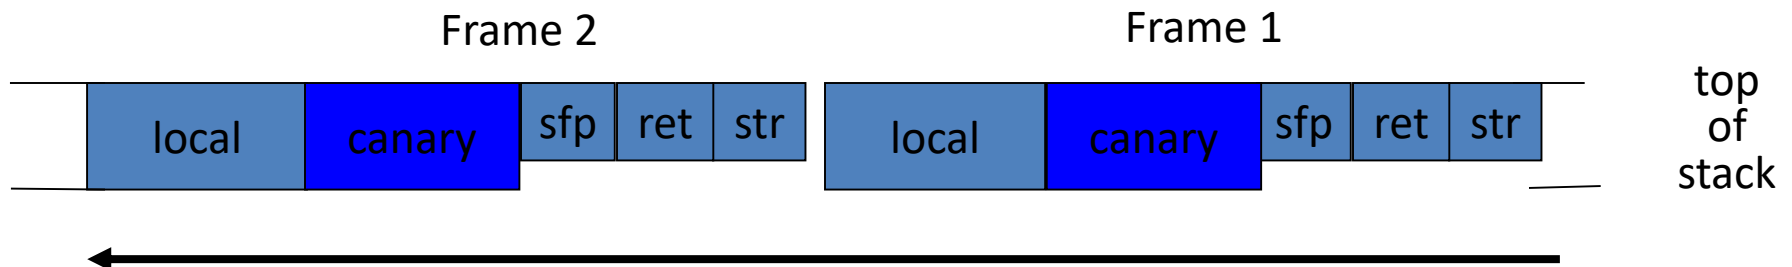
# Control Hijacking

---

## Run-time Defenses

# Run time checking: StackGuard

- Many run-time checking techniques ...
  - we only discuss methods relevant to overflow protection
- Solution 1: StackGuard
  - Run time tests for stack integrity.
  - Embed “canaries” in stack frames and verify their integrity prior to function return.



# Canary Types

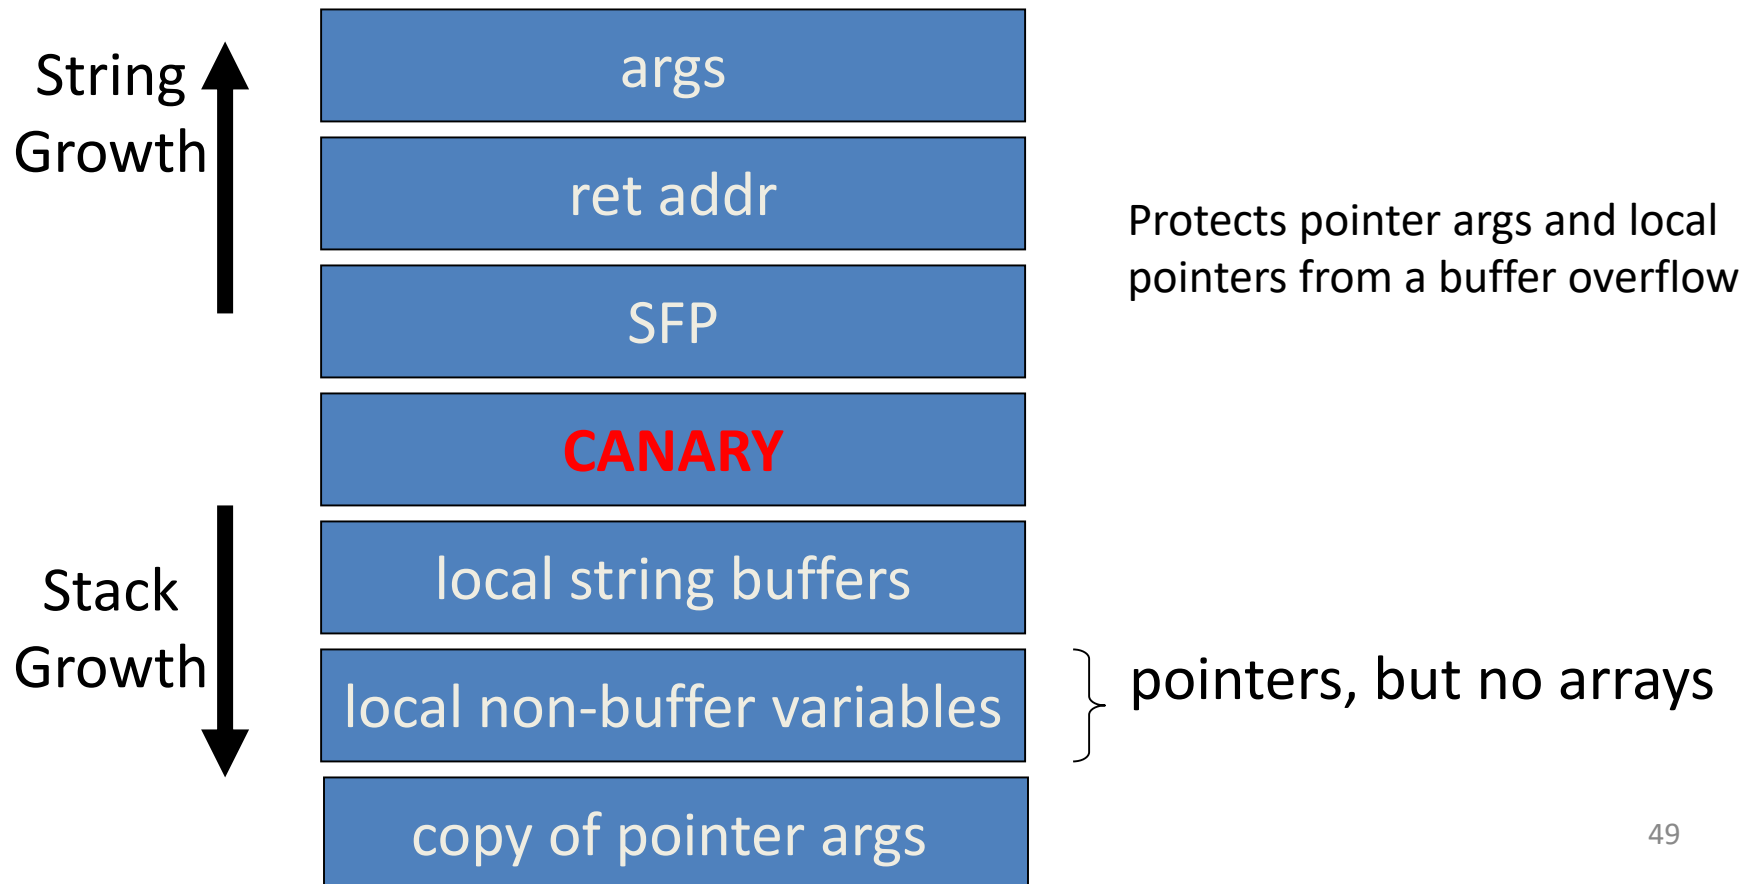
- Random canary:
  - Random string chosen at program startup.
  - Insert canary string into every stack frame.
  - Verify canary before returning from function.
    - Exit program if canary changed. Turns potential exploit into DoS.
  - To corrupt, attacker must learn current random string.
- Terminator canary:      Canary = {0, newline, linefeed, EOF}
  - String functions will not copy beyond terminator.
  - Attacker cannot use string functions to corrupt stack.

# StackGuard (Cont.)

- StackGuard implemented as a GCC patch
  - Program must be recompiled
- Minimal performance effects: 8% for Apache
- Note: Canaries do not provide full protection
  - Some stack smashing attacks leave canaries unchanged
- Heap protection: PointGuard
  - Protects function pointers and setjmp buffers by encrypting them: e.g. XOR with random cookie
  - Less effective, more noticeable performance effects

# StackGuard enhancements: ProPolice

- ProPolice (IBM) - gcc 3.4.1. (**-fstack-protector**)
  - Rearrange stack layout to prevent ptr overflow.





# MS Visual Studio /GS [since 2003]

Compiler /GS option:

- Combination of ProPolice and Random canary.
- If cookie mismatch, default behavior is to call **`_exit(3)`**

Function prolog:

```
sub esp, 8 // allocate 8 bytes for cookie
mov eax, DWORD PTR ___security_cookie
xor eax, esp // xor cookie with current esp
mov DWORD PTR [esp+8], eax // save in stack
```

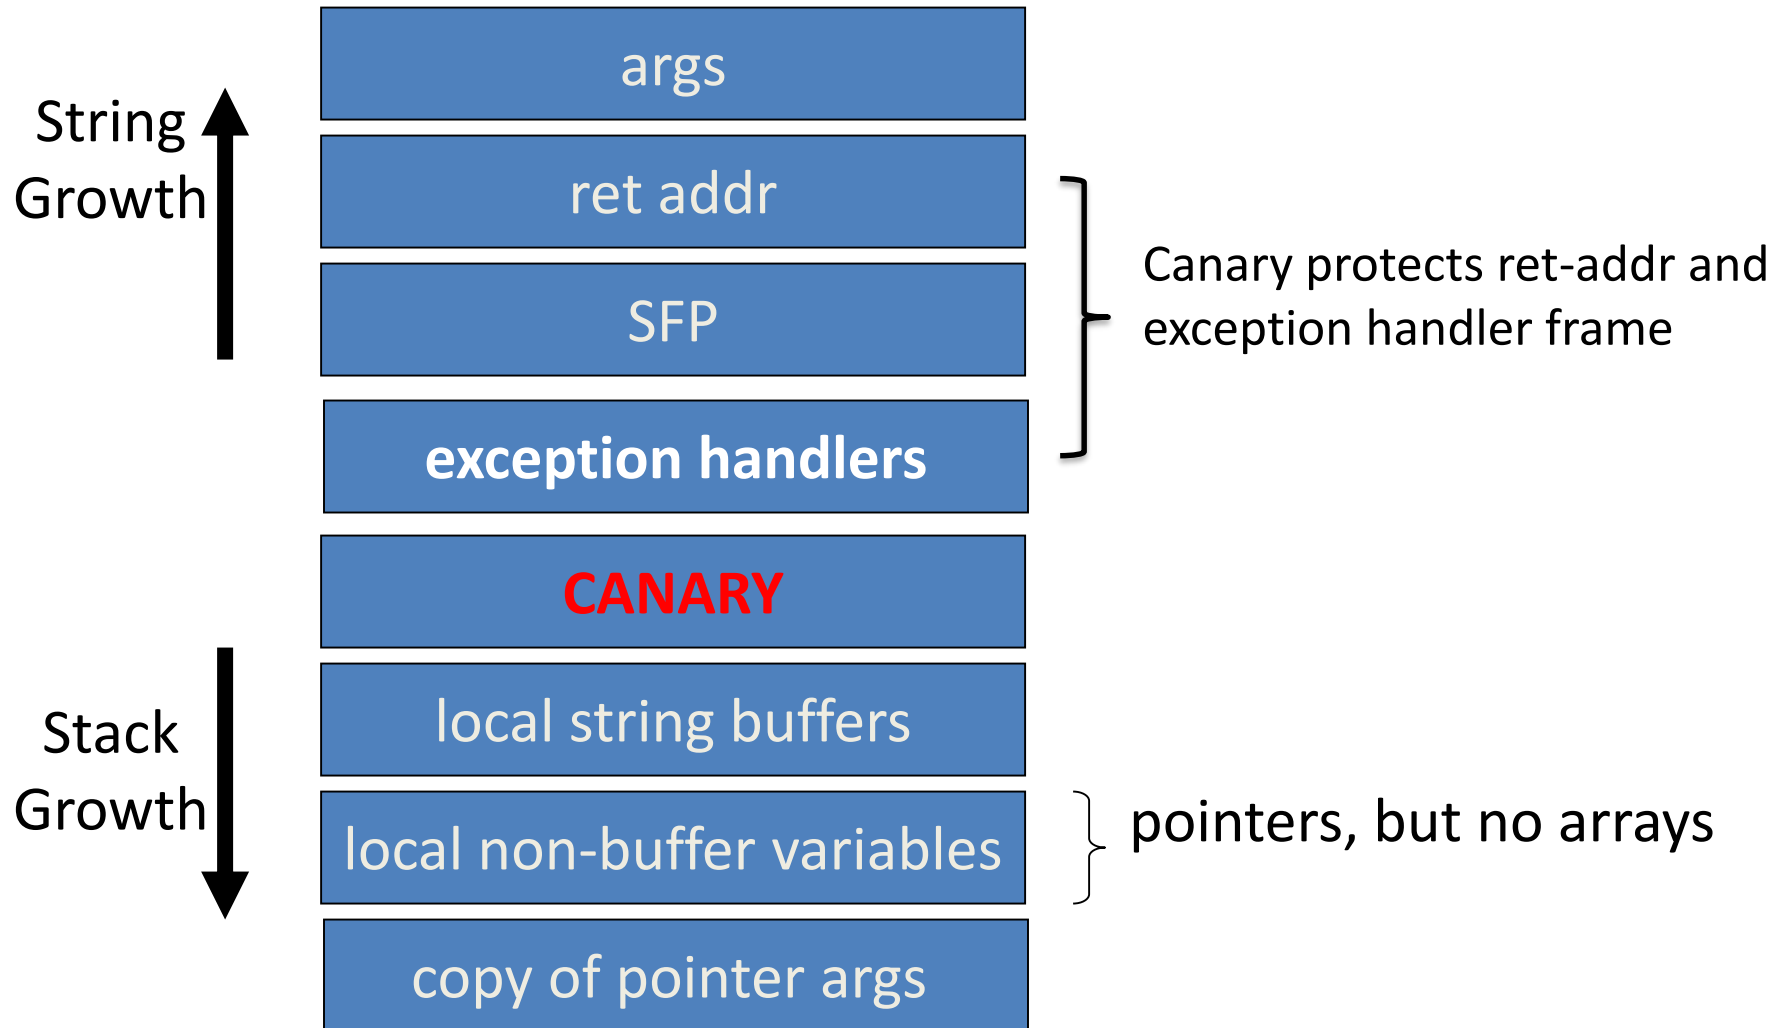
Function epilog:

```
mov ecx, DWORD PTR [esp+8]
xor ecx, esp
call @__security_check_cookie@4
add esp, 8
```

Enhanced /GS in Visual Studio 2010:

- /GS protection added to all functions, unless can be proven unnecessary

# /GS stack frame




# Evading /GS with exception handlers

- When exception is thrown, dispatcher walks up exception list until handler is found (else use default handler)

After overflow: handler points to attacker's code  
exception triggered  $\Rightarrow$  control hijack

Main point: exception is triggered before canary is checked



# Defenses: SAFESEH and SEHOP

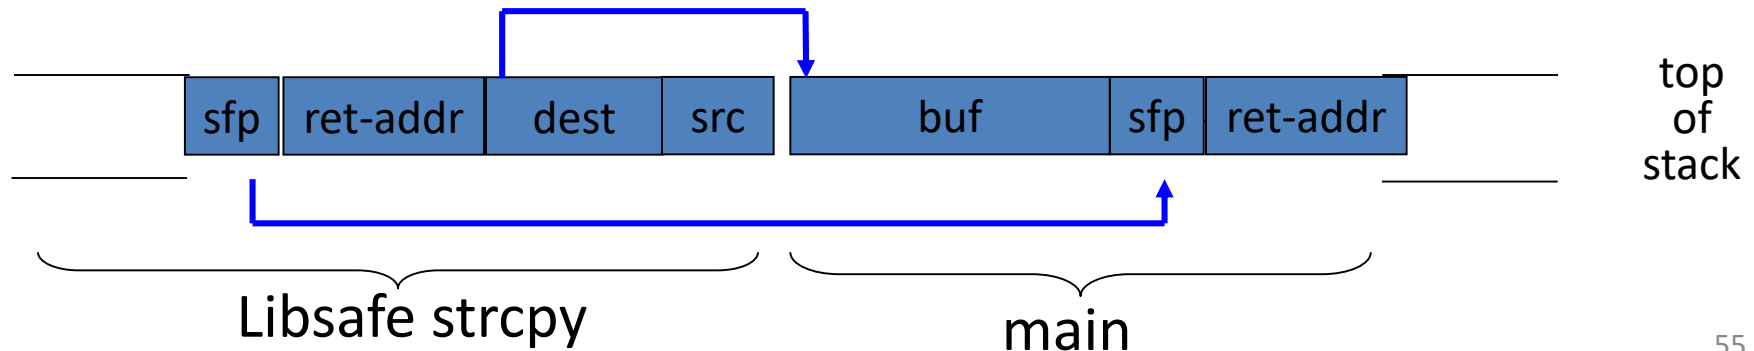
- **/SAFESSEH:** linker flag
  - Linker produces a binary with a table of safe exception handlers
  - System will not jump to exception handler not on list
- **/SEHOP:** platform defense (since win vista SP1)
  - Observation: SEH attacks typically corrupt the “next” entry in SEH list.
  - SEHOP: add a dummy record at top of SEH list
  - When exception occurs, dispatcher walks up list and verifies dummy record is there. If not, terminates process.

# Summary: Canaries are not full proof

- Canaries are an important defense tool, but do not prevent all control hijacking attacks:
  - Heap-based attacks still possible
  - Integer overflow attacks still possible
  - /GS by itself does not prevent Exception Handling attacks  
(also need SAFESSEH and SEHOP)

# What if can't recompile: Libsafe

- Solution 2: Libsafe (Avaya Labs)
  - Dynamically loaded library (no need to recompile app.)
  - Intercepts calls to `strcpy(dest, src)`
    - Validates sufficient space in current stack frame:  
 **$|\text{frame-pointer} - \text{dest}| > \text{strlen}(\text{src})$**
    - If so, does `strcpy`. Otherwise, terminates application



# More methods ...

## ➤ StackShield

- At function prologue, copy return address RET and SFP to “safe” location (beginning of data segment)
- Upon return, check that RET and SFP is equal to copy.
- Implemented as assembler file processor (GCC)

## ➤ Control Flow Integrity (CFI)

- A combination of static and dynamic checking
  - Statically determine program control flow
  - Dynamically enforce control flow integrity



# Control Hijacking

---

## Advanced Hijacking Attacks




# Heap Spray Attacks


A reliable method for exploiting heap overflows

# Heap-based control hijacking

- Compiler generated function pointers (e.g. C++ code)




- Suppose `vtable` is on the heap next to a string object:




# Heap-based control hijacking

- Compiler generated function pointers (e.g. C++ code)




- After overflow of **buf** we have:



# A reliable exploit?

```
<SCRIPT language="text/javascript">  
  shellcode = unescape("%u4343%u4343%...");  
  overflow-string = unescape("%u2332%u4276%...");  
  cause-overflow( overflow-string );    // overflow buf[ ]  
</SCRIPT>
```

Problem: attacker does not know where browser places **shellcode** on the heap




# Heap Spraying

[SkyLined 2004]

Idea:

1. use Javascript to spray heap with shellcode (and NOP slides)
2. then point vtable ptr anywhere in spray area



# Javascript heap spraying

```
var nop = unescape("%u9090%u9090")
while (nop.length < 0x100000)  nop += nop


var shellcode = unescape("%u4343%u4343%...");

var x = new Array ()
for (i=0; i<1000; i++) {
    x[i] = nop + shellcode;
}
```

- Pointing func-ptr almost anywhere in heap will cause shellcode to execute.

# Vulnerable buffer placement

- Placing vulnerable `buf[256]` next to object O:
  - By sequence of Javascript allocations and frees make heap look as follows:



- Allocate vuln. buffer in Javascript and cause overflow
- Successfully used against a Safari PCRE overflow [DHM'08]

# Many heap spray exploits

Date	Browser	Description
11/2004	IE	IFRAME Tag BO
04/2005	IE	DHTML Objects Corruption
01/2005	IE	.ANI Remote Stack BO
07/2005	IE	javaprxxy.dll COM Object
03/2006	IE	createTextRang RE
09/2006	IE	VML Remote BO
03/2007	IE	ADODB Double Free
09/2006	IE	WebViewFolderIcon setSlice
09/2005	FF	0xAD Remote Heap BO
12/2005	FF	compareTo() RE
07/2006	FF	Navigator Object RE
07/2008	Safari	Quicktime Content-Type BO

[RLZ'08]


- Improvements: Heap Feng Shui [S'07]
  - Reliable heap exploits **on IE** without spraying
  - Gives attacker full control of IE heap from Javascript



# (partial) Defenses

- Protect heap function pointers (e.g. PointGuard)
- Better browser architecture:
  - Store JavaScript strings in a separate heap from browser heap

- OpenBSD heap overflow protection:



- Nozzle [RLZ'08] : detect sprays by prevalence of code on heap

# References on heap spraying

- [1] **Heap Feng Shui in Javascript,**  
by A. Sotirov, *Blackhat Europe* 2007
- [2] **Engineering Heap Overflow Exploits with JavaScript**  
M. Daniel, J. Honoroff, and C. Miller, *WooT* 2008
- [3] **Nozzle: A Defense Against Heap-spraying Code Injection Attacks,**  
by P. Ratanaworabhan, B. Livshits, and B. Zorn
- [4] **Interpreter Exploitation: Pointer inference and JiT spraying,**  
by Dion Blazakis