

# Non-linear Models-IV

CS771: Introduction to Machine Learning

Purushottam Kar



# Outline of discussion

- Kernel Approximation Methods
- PML with kernels
- Neural networks

# Kernel methods can be slow 😞

- Need to work with indirect “dual” representations
- Although finite, these representations blow up with data size
- Prediction requires a full pass over data i.e.  $O(dn)$  time
- Will see some techniques to remedy this

# The Tale of a Trio of Techniques


- **Post-processing techniques**: learn the kernel SVM (a bit costly), but then make the model cheaper to store and predict
- **Approximate training techniques**: directly learn a kernel SVM model that is cheap to store and predict
- **Kernel approximation techniques**: use a different kernel than the one you wanted to, so that the new kernel mimics the original one but always gives models that are cheap to store and predict
- Kernel approximation is the most successful of the three

# Post Processing Techniques

- Learn kernel SVM, support vectors  $\{x_{i_j}, \alpha_{i_j}\}$
- Find a *reduced set* of  $k \ll \tilde{n}$  support vectors  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$  e.g. by using k-means clustering on original support vectors
- Re-compute  $\alpha$  values for these reduced set support vectors e.g. by running SVM again on  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$
- Burges and Scholkopf, Improving the Speed and Accuracy of SVMs, NIPS 1996.
- Cossalter et al. Adaptive Kernel Approximation for Large-Scale Non-Linear SVM Prediction, ICML 2011.


# Post Processing Techniques

- Learn kernel SVM, support vectors  $\{x_{i_j}, \alpha_{i_j}\}$
- Find a *reduced set* of  $k \ll \tilde{n}$  support vectors  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$  e.g. by using k-means clustering on original support vectors
- Re-compute  $\alpha$  values for these reduced set support vectors e.g. by running SVM again on  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$
- Burges and Scholkopf, Improving the Speed and Accuracy of SVMs, NIPS 1996.
- Cossalter et al. Adaptive Kernel Approximation for Large-Scale Non-Linear SVM Prediction, ICML 2011.




# Post Processing Techniques

- Learn kernel SVM, support vectors  $\{x_{i_j}, \alpha_{i_j}\}$
- Find a *reduced set* of  $k \ll \tilde{n}$  support vectors  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$  e.g. by using k-means clustering on original support vectors
- Re-compute  $\alpha$  values for these reduced set support vectors e.g. by running SVM again on  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$
- Burges and Scholkopf, Improving the Speed and Accuracy of SVMs, NIPS 1996.
- Cossalter et al. Adaptive Kernel Approximation for Large-Scale Non-Linear SVM Prediction, ICML 2011.



# Post Processing Techniques


- Learn kernel SVM, support vectors  $\{x_{i_j}, \alpha_{i_j}\}$
- Find a *reduced set* of  $k \ll \tilde{n}$  support vectors  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$  e.g. by using k-means clustering on original support vectors
- Re-compute  $\alpha$  values for these reduced set support vectors e.g. by running SVM again on  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$
- Burges and Scholkopf, Improving the Speed and Accuracy of SVMs, NIPS 1996.
- Cossalter et al. Adaptive Kernel Approximation for Large-Scale Non-Linear SVM Prediction, ICML 2011.






# Post Processing Techniques

- Learn kernel SVM, support vectors  $\{x_{i_j}, \alpha_{i_j}\}$
- Find a *reduced set* of  $k \ll \tilde{n}$  support vectors  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$  e.g. by using k-means clustering on original support vectors
- Re-compute  $\alpha$  values for these reduced set support vectors e.g. by running SVM again on  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$
- Burges and Scholkopf, Improving the Speed and Accuracy of SVMs, NIPS 1996.
- Cossalter et al. Adaptive Kernel Approximation for Large-Scale Non-Linear SVM Prediction, ICML 2011.




# Post Processing Techniques

- Learn kernel SVM, support vectors  $\{x_{i_j}, \alpha_{i_j}\}$
- Find a *reduced set* of  $k \ll \tilde{n}$  support vectors  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$  e.g. by using k-means clustering on original support vectors
- Re-compute  $\alpha$  values for these reduced set support vectors e.g. by running SVM again on  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$
- Burges and Scholkopf, Improving the Speed and Accuracy of SVMs, NIPS 1996.
- Cossalter et al. Adaptive Kernel Approximation for Large-Scale Non-Linear SVM Prediction, ICML 2011.




# Post Processing Techniques

- Learn kernel SVM, support vectors  $\{x_{i_j}, \alpha_{i_j}\}$
- Find a *reduced set* of  $k \ll \tilde{n}$  support vectors  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$  e.g. by using k-means clustering on original support vectors
- Re-compute  $\alpha$  values for these reduced set support vectors e.g. by running SVM again on  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$
- Burges and Scholkopf, Improving the Speed and Accuracy of SVMs, NIPS 1996.
- Cossalter et al. Adaptive Kernel Approximation for Large-Scale Non-Linear SVM Prediction, ICML 2011.




# Post Processing Techniques

- Learn kernel SVM, support vectors  $\{x_{i_j}, \alpha_{i_j}\}$
- Find a *reduced set* of  $k \ll \tilde{n}$  support vectors  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$  e.g. by using k-means clustering on original support vectors
- Re-compute  $\alpha$  values for these reduced set support vectors e.g. by running SVM again on  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$
- Burges and Scholkopf, Improving the Speed and Accuracy of SVMs, NIPS 1996.
- Cossalter et al. Adaptive Kernel Approximation for Large-Scale Non-Linear SVM Prediction, ICML 2011.



# Post Processing Techniques

- Learn kernel SVM, support vectors  $\{x_{i_j}, \alpha_{i_j}\}$
- Find a *reduced set* of  $k \ll \tilde{n}$  support vectors  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$  e.g. by using k-means clustering on original support vectors
- Re-compute  $\alpha$  values for these reduced set support vectors e.g. by running SVM again on  $\{\tilde{x}_{i_1}, \dots, \tilde{x}_{i_k}\}$
- Burges and Scholkopf, Improving the Speed and Accuracy of SVMs, NIPS 1996.
- Cossalter et al. Adaptive Kernel Approximation for Large-Scale Non-Linear SVM Prediction, ICML 2011.



# Approximate Training Techniques

- Notice that support vectors are always a subset of training data
- Maybe removing this restriction can reduce their number?
- Learn support vectors as well (not necessarily training points)!
- Learn vectors  $\mathbf{z}^1, \dots, \mathbf{z}^k \in \mathbb{R}^d$  and weights  $\alpha_1, \dots, \alpha_k \in \mathbb{R}$  so that

$$\mathbf{w} = \sum_{i=1}^k \alpha_i \cdot \phi_K(\mathbf{z}^i)$$

is a good model (classifier, regressor etc)

- $k$  chosen based on budget (space, time) of application
- $\mathcal{O}(kd)$  storage and  $\mathcal{O}(kd)$  time for prediction
- Joachims and Yu. Sparse Kernel SVMs via Cutting-Plane Training, Machine Learning 76(2):179-193, 2009
- Tsang et al. Core Vector Machines, JMLR 6:363-392, 2005





# Approximate Training Techniques

- Notice that support vectors are always a subset of training data
- Maybe removing this restriction can reduce the number of support vectors?
- Learn support vectors as well (not necessarily training points)!
- Learn vectors  $\mathbf{z}^1, \dots, \mathbf{z}^k \in \mathbb{R}^d$  and weights  $\alpha_1, \dots, \alpha_k \in \mathbb{R}$  so that

$$\mathbf{w} = \sum_{i=1}^k \alpha_i \cdot \phi_K(\mathbf{z}^i)$$

is a good model (classifier, regressor etc)

- $k$  chosen based on budget (space, time) of application
- $\mathcal{O}(kd)$  storage and  $\mathcal{O}(kd)$  time for prediction
- Joachims and Yu. Sparse Kernel SVMs via Cutting-Plane Training, Machine Learning 76(2):179-193, 2009
- Tsang et al. Core Vector Machines, JMLR 6:363-392, 2005

$\phi_K$  is the map  
for kernel  $K$



# Kernel Approximation

Oct 18, 2017





# Landmarking

- Given: training set  $S = \{x^1, \dots, x^n\}$  and kernel  $K$
- Select  $k \ll n$  landmarks
$$\hat{S} = \{\hat{x}^1, \dots, \hat{x}^k\} \subset S$$
- Choice may be random or careful (more expensive)
- Use landmarks to create a new  $k$ -dim. feature representation
$$\hat{\phi}(x) = [K(x, \hat{x}^1), \dots, K(x, \hat{x}^k)]$$
- Now use  $\hat{\phi}(x)$  to perform classification, regression, etc
- Can be theoretically shown that if  $K$  was nice, so will be  $\hat{K}$
- No agony of high dim-feature map with  $\hat{K}$
- Balcan and Blum. On a Theory of Learning with Similarity Functions, ICML 2006.
- K. and Jain. Similarity-based Learning via Data driven Embeddings, NIPS 2011.

# Landmarking

- Given training set  $S = \{x^1, \dots, x^n\}$  and kernel  $K$

Since  $k$  is chosen to be small, so use linear SVM/RR over  $\hat{\phi}(x)$  directly

$$\hat{S} = \{\hat{x}^1, \dots, \hat{x}^k\} \subset S$$

Can think of  $\hat{\phi}$  as giving us a new kernel  $\hat{K}(x, y) = \langle \hat{\phi}(x), \hat{\phi}(y) \rangle$

- Choice may be random or careful (more expensive)

- Use landmarks to create a new  $k$ -dim. feature representation

$$\hat{\phi}(x) = [K(x, \hat{x}^1), \dots, K(x, \hat{x}^k)]$$

$\mathcal{O}(kd)$  model size and  $\mathcal{O}(kd)$  prediction time

- Now use  $\hat{\phi}(x)$  to perform classification, regression, etc.

- Can be theoretically shown that if  $K$  was nice, so will be  $\hat{K}$

- No agony of high dim-feature map with  $\hat{K}$

Work with non-Mercer kernels too!

- Balcan and Blum. On a Theory of Learning with Similarity Functions, ICM 2004.
- K. and Jain. Similarity-based Learning via Data driven Embeddings, NIPS 2011.

# Nystrom Method

- A more careful implementation of landmarking
- **Basic idea:** landmarks may be correlated – decorrelate them
- Recall landmark set  $\hat{S}$  of size  $k$  gave us a map  $\hat{\phi}$  that maps to  $\mathbb{R}^k$
- Let  $\hat{G} \in \mathbb{R}^{k \times k}$  be Gram matrix over landmark set  $\hat{S}$  and let its eigendecomposition be  $\hat{G} = U\Lambda U^\top$  where  $U = [u^1, \dots, u^k] \in \mathbb{R}^{k \times k}$  is the matrix of eigenvectors and  $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_k)$  be eigenvalues
- Nystrom method defines the similarity between  $x, y$  as

$$\hat{\phi}(x)^\top G^\dagger \hat{\phi}(y)$$

- Nystrom features are modified version of landmarking feature  $\hat{\phi}$

$$\tilde{\phi}(x) = \sqrt{\Lambda^{-1}} U^\top \hat{\phi}(x)$$

if any  $\lambda_i = 0$ , remove that eigenvalue from  $\Lambda$  and vector from  $U$

# Nystrom Method

- A more careful implementation of landmarking
- **Basic idea:** landmarks may be correlated – decorrelate them
- Recall landmark set  $\hat{S}$  of size  $k$  gave us a map  $\hat{\phi}$  that maps to  $\mathbb{R}^k$
- Let  $\hat{G} \in \mathbb{R}^{k \times k}$  be Gram matrix over landmark set  $\hat{S}$  and let its eigendecomposition be  $\hat{G} = U\Lambda U^\top$  where  $U = [u^1, \dots, u^k] \in \mathbb{R}^{k \times k}$  is the matrix of eigenvectors and  $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_k)$  is the diagonal matrix of eigenvalues. **Pseudoinverse**
- Nystrom method defines the similarity between  $x, y$  as

$$\hat{\phi}(x)^\top G^\dagger \hat{\phi}(y)$$

- Nystrom features are modified version of landmarking feature  $\hat{\phi}$

$$\tilde{\phi}(x) = \sqrt{\Lambda^{-1}} U^\top \hat{\phi}(x)$$

if any  $\lambda_i = 0$ , remove that eigenvalue from  $\Lambda$  and vector from  $U$

# Nystrom Method

- The Nystrom feature also gives us a new kernel  $\tilde{K}$ 
$$\tilde{K}(x, y) = \tilde{\phi}(x)^\top \tilde{\phi}(y) = \hat{\phi}(x)^\top U \Lambda^{-1} U^\top \hat{\phi}(y)$$
- Note that the Gram matrices corresponding to the landmarking kernels  $\hat{K}$  as well as Nystrom kernel  $\tilde{K}$  are always rank atmost  $k$
- Interesting note: suppose actual kernel  $K$  has map  $\phi_K(x) \in \mathbb{R}^D$
- Let  $\Phi_{\hat{S}} = [\phi(\hat{x}^i)]_{i=1, \dots, k} \in \mathbb{R}^{D \times k}$  where  $\hat{S} = \{\hat{x}^1, \dots, \hat{x}^k\}$  is landmark set
- This means  $\hat{\phi}(x) = \Phi_{\hat{S}}^\top \phi_K(x)$  and  $\hat{G} = \Phi_{\hat{S}}^\top \Phi_{\hat{S}}$  i.e.
$$\tilde{K}(x, y) = \phi_K(x)^\top \Phi_{\hat{S}} (\Phi_{\hat{S}}^\top \Phi_{\hat{S}})^\dagger \Phi_{\hat{S}}^\top \phi_K(y)$$
- Takes more time  $O(k^2 + kd)$  to construct Nystrom feature map
- Williams and Seeger. Using the Nystrom Method to Speed Up Kernel Machines, NIPS 2000
- Yang et al. Nystrom Method vs Random Fourier Features, NIPS 2012

# Nystrom Method

- The Nystrom feature also gives us a new kernel  $\tilde{K}$ 
$$\tilde{K}(x, y) = \tilde{\phi}(x)^\top \tilde{\phi}(y) = \hat{\phi}(x)^\top U \Lambda^{-1} U^\top \hat{\phi}(y)$$
- Note that the Gram matrices corresponding to the landmarking kernels  $\hat{K}$  as well as Nystrom kernel  $\tilde{K}$  are always rank atmost  $k$
- Interesting note: suppose actual kernel  $K$  has map  $\phi_K(x) \in \mathbb{R}^D$
- Let  $\Phi_{\hat{S}} = [\phi(\hat{x}^i)]_{i=1, \dots, k} \in \mathbb{R}^{D \times k}$  where  $\hat{S} = \{\hat{x}^1, \dots, \hat{x}^k\}$  is landmark set
- This means  $\hat{\phi}(x) = \Phi_{\hat{S}}^\top \phi_K(x)$  and  $\hat{G} = \Phi_{\hat{S}}^\top \Phi_{\hat{S}}$  i.e. Decorrelation
- $$\tilde{K}(x, y) = \phi_K(x)^\top \Phi_{\hat{S}} (\Phi_{\hat{S}}^\top \Phi_{\hat{S}})^\dagger \Phi_{\hat{S}}^\top \phi_K(y)$$
- Takes more time  $O(k^2 + kd)$  to construct Nystrom feature map
- Williams and Seeger. Using the Nystrom Method to Speed Up Kernel Machines, NIPS 2000
- Yang et al. Nystrom Method vs Random Fourier Features, NIPS 2012

# Explicit Feature Constructions

- Realize that high dim. of feature map  $\phi_K$  is root of all problems
- If  $\phi_K$  were small dim. then training, storage, testing much easier
- Given a Mercer kernel  $K: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  can we construct  $\bar{\phi}: \mathcal{X} \rightarrow \mathbb{R}^k$ 
  - $k$  should not be too large so we can use  $\bar{\phi}$  explicitly
  - It should be easy to map  $x \mapsto \bar{\phi}(x)$
  - $\bar{\phi}$  should act as an approx feature map for  $K$  i.e. for all  $x, y \in \mathcal{X}$ 
$$\langle \bar{\phi}(x), \bar{\phi}(y) \rangle =: \bar{K}(x, y) \approx K(x, y) = \langle \phi_K(x), \phi_K(y) \rangle$$
- Note that landmarking and Nystrom do not seek to ensure that  $\hat{K}$  or  $\tilde{K}$  values approximate  $K$  but  $\bar{K}$  should approximate  $K$  values
- Why should such  $\bar{\phi}$  even exist?

# Random Feature Constructions

- Several popular Mercer kernels have a peculiar form

$$K(x, y) = \mathbb{E}_{\omega \sim \mathcal{D}_K} [K_\omega(x, y)]$$

- $\omega$  is an auxiliary variable (depending on the kernel,  $\omega \in \mathbb{N}, \mathbb{R}, \mathbb{R}^d$ )
- $\mathcal{D}_K$  is a distribution that depends on kernel  $K$  and known to us
- $K_\omega$  is a very “simple” Mercer kernel, has a one-dim. feature map

$$K_\omega(x, y) = \langle \phi_\omega(x), \phi_\omega(y) \rangle$$

$$\phi_\omega: \mathcal{X} \rightarrow \mathbb{R}$$

- Sample several  $\omega_1, \dots, \omega_k$  and define the map

$$\bar{\phi}: x \mapsto [\phi_{\omega_1}(x), \dots, \phi_{\omega_k}(x)] \in \mathbb{R}^k$$

- Can theoretically prove that with high probability

$$\langle \bar{\phi}(x), \bar{\phi}(y) \rangle \approx K(x, y)$$



# Random Feature Constructions

- Gaussian/Laplacian kernels

$$K(\mathbf{x}, \mathbf{y}) = \mathbb{E}_{\omega \sim \mathcal{D}_K} [\cos(\omega^\top \mathbf{x}) \cos(\omega^\top \mathbf{y})]$$
$$\phi_\omega: \mathbf{x} \mapsto \cos(\omega^\top \mathbf{x})$$

Note that auxiliary variable is a vector here  $\omega \in \mathbb{R}^d$

Rahimi and Recht, Random Features for Large Scale Kernel Machines, NIPS 2007

- Intersection kernel

Maji and Berg, Max-margin Additive Classifiers for Detect, ICCV 2009.

- Homogeneous kernels

Vedaldi and Zisserman. Efficient Additive Kernels via Explicit Feature Maps, CVPR 2010

- Polynomial kernels

K. and Karnick. Random Feature Maps for Dot Product Kernels. AISTATS 2012

# Other kernel approximation approaches

- Use decision trees to compute similarity between two points and use that as kernel – extremely fast prediction

Jose et al. Local Deep Kernel Learning, ICML 2013.

- Learn these kernel approximations in a task-dependent manner

Perronnin et al. d Yan Liu. Large-scale Image Categorization with Explicit Data embedding, CVPR 2010.

# PML with Kernels

## Gaussian Processes

Oct 18, 2017



# Priors and Posteriors

- How can we argue about priors and posteriors in an RKHS  $\mathcal{H}_K$ ?
- Details too advanced (covered in CS772, CS775, CS698X)
- Basic idea: argue about distributions over functions  $f: \mathcal{X} \rightarrow \mathbb{R}$
- Gaussian processes is one such family of distributions
$$f \sim \text{GP}(\mu, K)$$
$$\mu: \mathcal{X} \rightarrow \mathbb{R} \text{ is the } \textit{mean} \text{ function and } \kappa \text{ is the covariance kernel}$$
- What does it mean to sample a function?
- Think of sampling a very very long vector (imprecise way though)
- Let  $|\mathcal{X}| = N < \infty$  with  $\mathcal{X} = \{x^1, x^2, \dots, x^N\}$
- Then can think of  $f: \mathcal{X} \rightarrow \mathbb{R}$  as a vector in  $\mathbb{R}^N$ 
$$f = [f(x^1), \dots, f(x^N)]$$

# Gaussian Processes

- For  $|\mathcal{X}| = N < \infty$ , we say a function  $f$  is sampled from  $\text{GP}(\mu, K)$  if
$$f \sim \mathcal{N}(\mu, G)$$
where  $\mu \in \mathbb{R}^N$  is mean fn. and  $G \in \mathbb{R}^{N \times N}$  with  $G_{ij} = K(x^i, x^j)$
- Note that  $f$  need not be linear etc, can be very complex
- Gaussian processes popularly use a Gaussian kernel for  $K$
- Note that the Gaussian kernel  $K$  forces  $f$  to be *smooth* i.e. if two points  $x^i, x^j \in \mathcal{X}$  are close i.e.  $\|x^i - x^j\|_2$  is small then functions  $f$  that take very different values on these points get low prob.
- **Exercise:** verify this yourself
- Mean function is taken to be zero (unless we have other reasons)

# Gaussian Process Regression

- Solve a regression problem  $\{x^i, y^i\}_{i=1, \dots, n}$ ,  $x^i \in \mathcal{X}$ ,  $y^i \in \mathbb{R}$  and  $n \ll N$
- Prior dist. (GP)  $f \sim \text{GP}(0, K)$
- Likelihood dist. (Gaussian)  $y^i | f \sim \mathcal{N}(f(x^i), \sigma^2)$
- Note: GP makes sense even if  $\mathcal{X}$  is set of vectors, images, text etc
- Can do regression over vectors, images as we did in kernel RR
- Let  $\mathbf{y} = [y^1, \dots, y^n]^\top \in \mathbb{R}^n$
- Using a very special property of Gaussians we can show
$$\mathbf{y} \sim \mathcal{N}(\mathbf{0}, G_n + \sigma^2 \cdot I_n)$$
where  $\mathbf{0} \in \mathbb{R}^n$  and  $G_n \in \mathbb{R}^{n \times n}$  is the Gram matrix of training data

# Gaussian Process Regression

- Solve  $\mathbf{v} \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma)$
  - Prior  $\mathbf{v}_S \sim \mathcal{N}(\boldsymbol{\mu}_S, \Sigma_{S,S})$
  - Likelihood  $\Sigma_{S,S} \in \mathbb{R}^{|S| \times |S|}$
- If a vector  $\mathbf{v} \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma)$  then  $\mathbf{v}_S \sim \mathcal{N}(\boldsymbol{\mu}_S, \Sigma_{S,S})$
- If a vector  $\mathbf{v} \in \mathbb{R}^n$  is distributed according to a Gaussian, then for every subset  $S \subset [n]$ , the sub-vector  $\mathbf{v}_S = [\mathbf{v}_i]_{i \in S} \in \mathbb{R}^{|S|}$  is also a Gaussian vector!
- $\mathbf{y}$  is just a subvector of  $f$
- Let  $\mathbf{y} = [y^1, \dots, y^n]^T \in \mathbb{R}^n$
  - Using a very special property of Gaussians we can show  $\mathbf{y} \sim \mathcal{N}(\mathbf{0}, G_n + \sigma^2 \cdot I_n)$
- where  $\mathbf{0} \in \mathbb{R}^n$  and  $G_n \in \mathbb{R}^{n \times n}$  is the Gram matrix of training data

# Gaussian Process Regression

- So we have  $\mathbf{y} \sim \mathcal{N}(\mathbf{0}, G_n + \sigma^2 \cdot I_n)$
- Now a new test point comes along  $\tilde{x} \in \mathcal{X}$ . Can we predict  $\tilde{y}$ ?
- Let  $\tilde{\mathbf{y}} = [\mathbf{y}, \tilde{y}] \in \mathbb{R}^{n+1}$ ,  $G_{n+1}$  be the Gram matrix over  $\{x^i\}_{i=1,\dots,n} \cup \tilde{x}$
- Previous slide gives us  $\tilde{\mathbf{y}} \sim \mathcal{N}(\mathbf{0}, G_{n+1} + \sigma^2 \cdot I_{n+1})$
- Let  $\tilde{\mathbf{g}} = [K(x^1, \tilde{x}), \dots, K(x^n, \tilde{x})]^\top \in \mathbb{R}^n$
- Then we can show that

$$\mathbb{P} \left[ \tilde{y} \mid \tilde{x}, \{x^i, y^i\}_{i=1,\dots,n} \right] = \mathcal{N}(\tilde{\mu}, \tilde{\sigma}^2)$$

$$\tilde{\mu} = \tilde{\mathbf{g}}^\top (G_n + \sigma^2 \cdot I_n)^{-1} \mathbf{y}$$

$$\tilde{\sigma}^2 = K(\tilde{x}, \tilde{x}) + \sigma^2 - \tilde{\mathbf{g}}^\top (G_n + \sigma^2 \cdot I_n)^{-1} \tilde{\mathbf{g}}$$



# Gaussian Process Regression

- So we have  $\mathbf{y} \sim \mathcal{N}(\mathbf{0}, G_n + \sigma^2 \cdot I_n)$
- Now a new test point comes along  $\tilde{x} \in \mathcal{X}$ . Can we predict  $\tilde{y}$ ?
- Let  $\tilde{\mathbf{y}} = [\mathbf{y}, \tilde{y}] \in \mathbb{R}^{n+1}$ ,  $G_{n+1}$  be the Gram matrix over  $(x^1, \dots, x^n, \tilde{x})$
- Previous slide gives us  $\tilde{\mathbf{y}} \sim \mathcal{N}(\mathbf{0}, G_{n+1} + \sigma^2 \cdot I_{n+1})$
- Let  $\tilde{\mathbf{g}} = [K(x^1, \tilde{x}), \dots, K(x^n, \tilde{x})]^\top \in \mathbb{R}^n$
- Then we can show that

$$\mathbb{P}[\tilde{y} \mid \tilde{x}, \{x^i, y^i\}_{i=1, \dots, n}] = \mathcal{N}(\tilde{\mu}, \tilde{\sigma}^2)$$

$$\tilde{\mu} = \tilde{\mathbf{g}}^\top (G_n + \sigma^2 \cdot I_n)^{-1} \mathbf{y}$$

$$\tilde{\sigma}^2 = K(\tilde{x}, \tilde{x}) + \sigma^2 - \tilde{\mathbf{g}}^\top (G_n + \sigma^2 \cdot I_n)^{-1} \tilde{\mathbf{g}}$$

Predictive posterior

Verify that the mean  $\tilde{\mu}$  is nothing but the kernel RR solution!

# A few thoughts

- GP regression is a Bayesian counterpart to kernel RR
- Similar cost for storing model, making predictions
- GP gives additional information about variance in prediction just as Bayesian models usually do (ref. Bayesian linear regression)
- Can apply accelerated learning techniques to GPs as well
- Can use GPs to perform kernel dim-redu as well
- Just as we did online MAP, can do online GP as well
- Btw, can do online kernel SVM, online kernel RR as well 😊
- Kernel perceptron is already online

# Neural Networks

Oct 18, 2017



35

CS771: Intro to ML

# Disclaimers

- Field is progressing rapidly – newer methods being proposed
- Some of the mentors, even some course students, more experienced with neural networks than the instructor
- Will cover very basics and essentials

# Back to Kernels first

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$
- A linear model over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$

# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$
- A linear model over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$

# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$
- A linear model over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$

$\mathbf{x}_1$

$\mathbf{x}_2$

# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$
- A linear model over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$

1

$\mathbf{x}_1$


$\mathbf{x}_2$



# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$


- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$
- A linear model over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$



# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$


- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$
- A linear model over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$



# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$


- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$
- A linear model over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$



# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$


- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$
- A linear model over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$



# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$

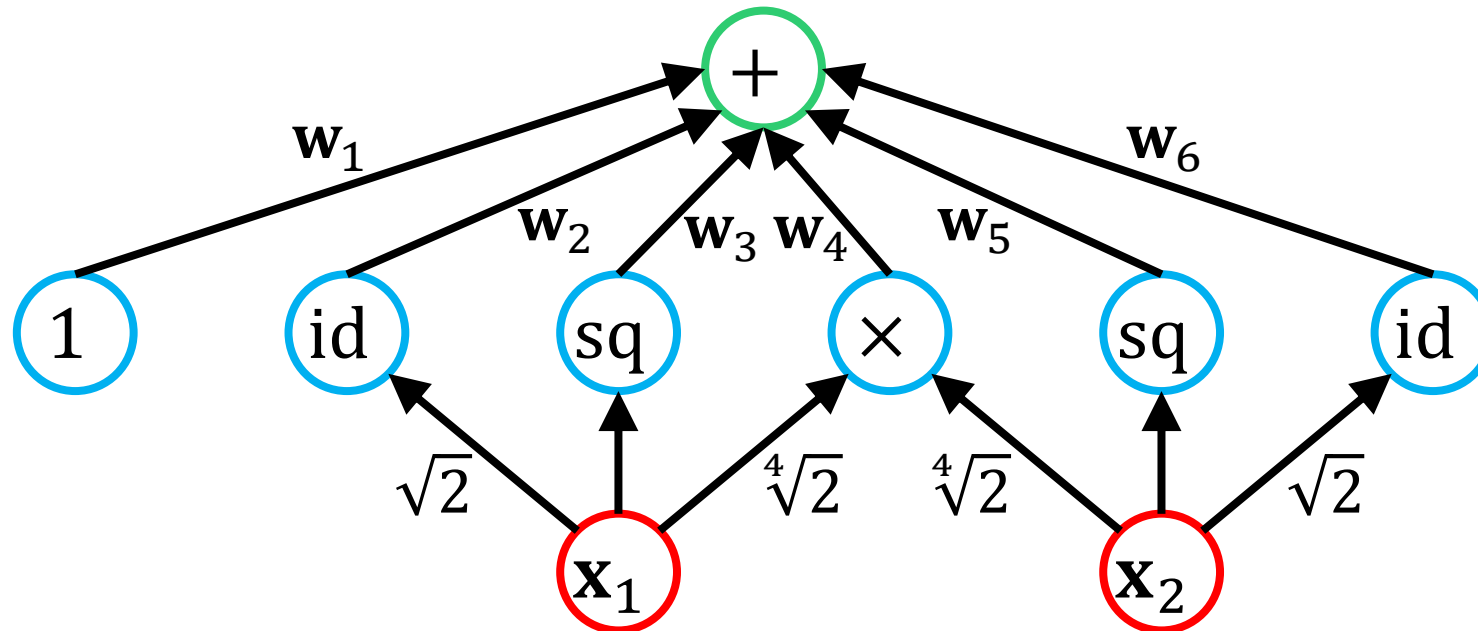
- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$
- A linear model over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$



# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$

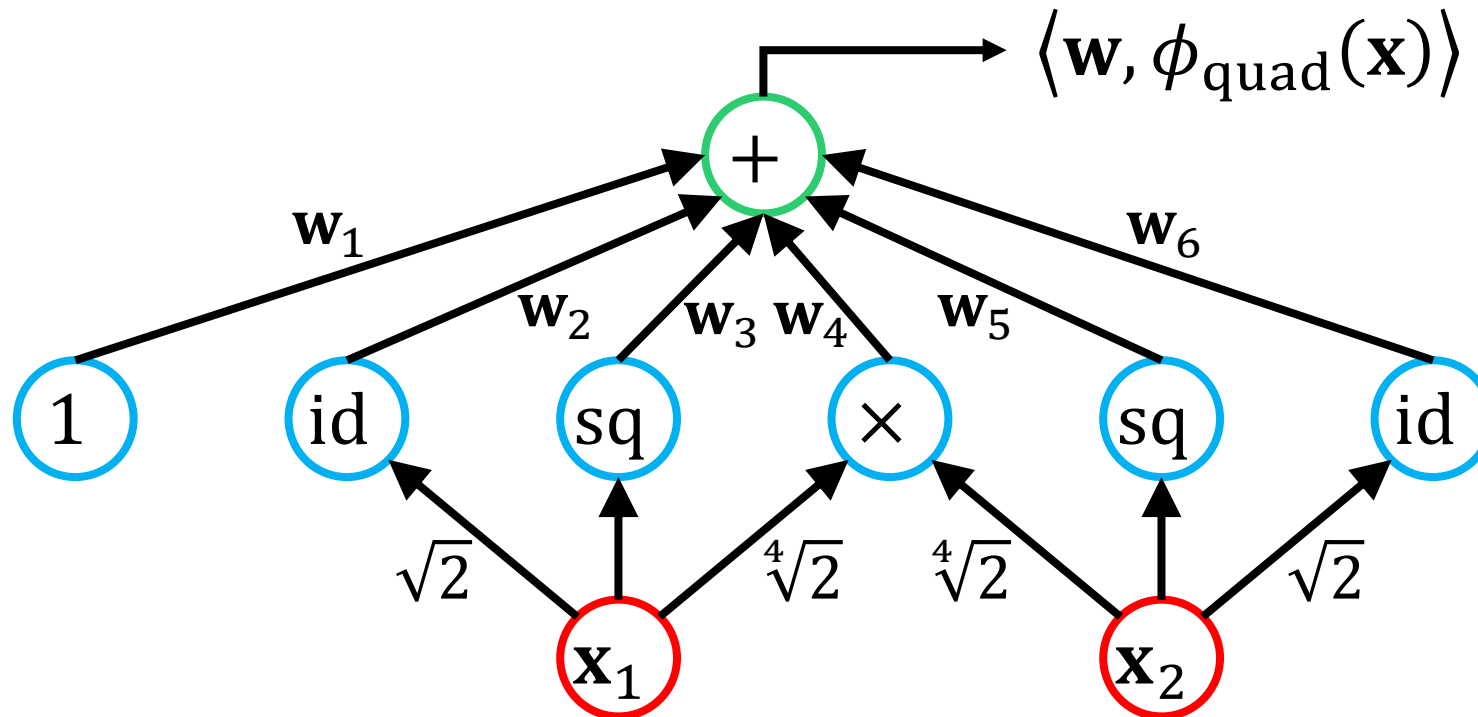
- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$
- A linear model over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$



# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$
- A linear model over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$



# Back to Kernels first

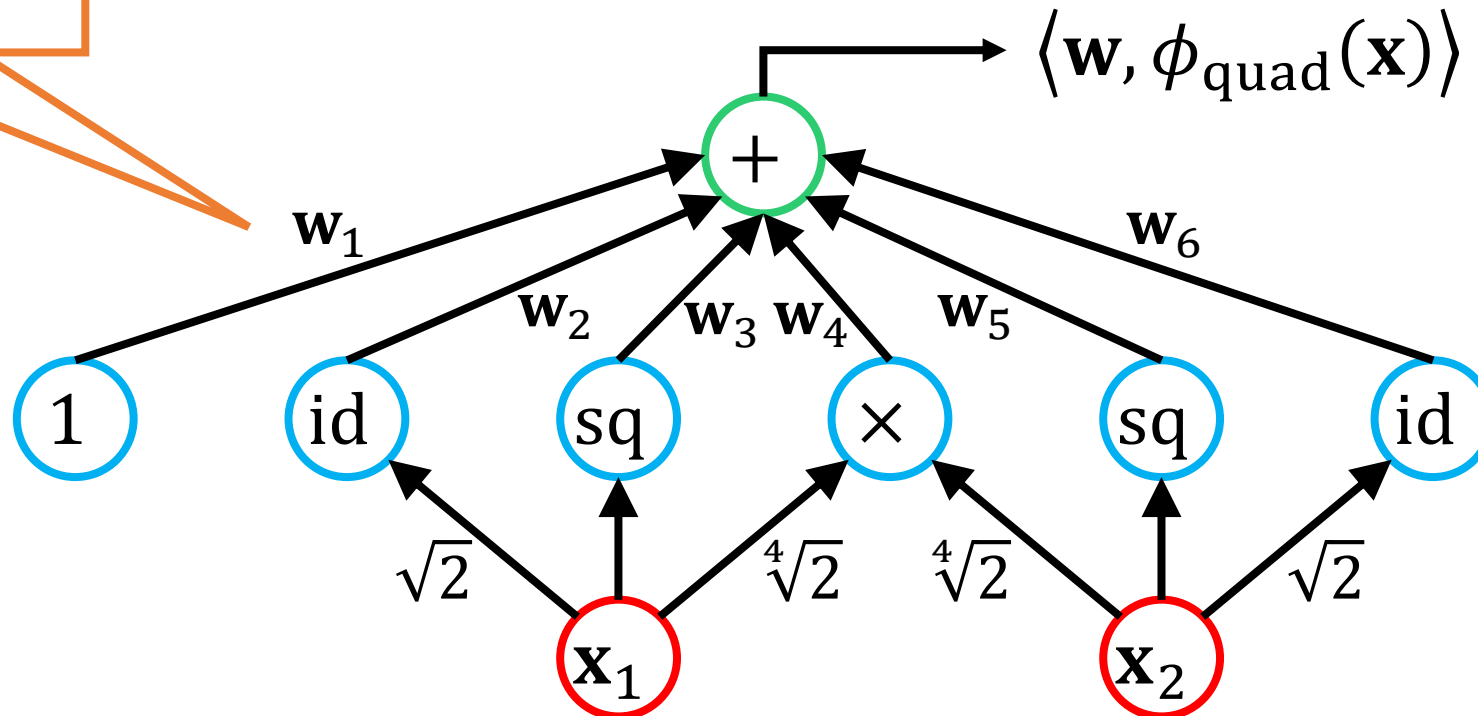
Can represent any quadratic fn over  $\mathbf{x}$

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$

- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$

$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$

Training an SVM using GD/CD tunes these weights over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$





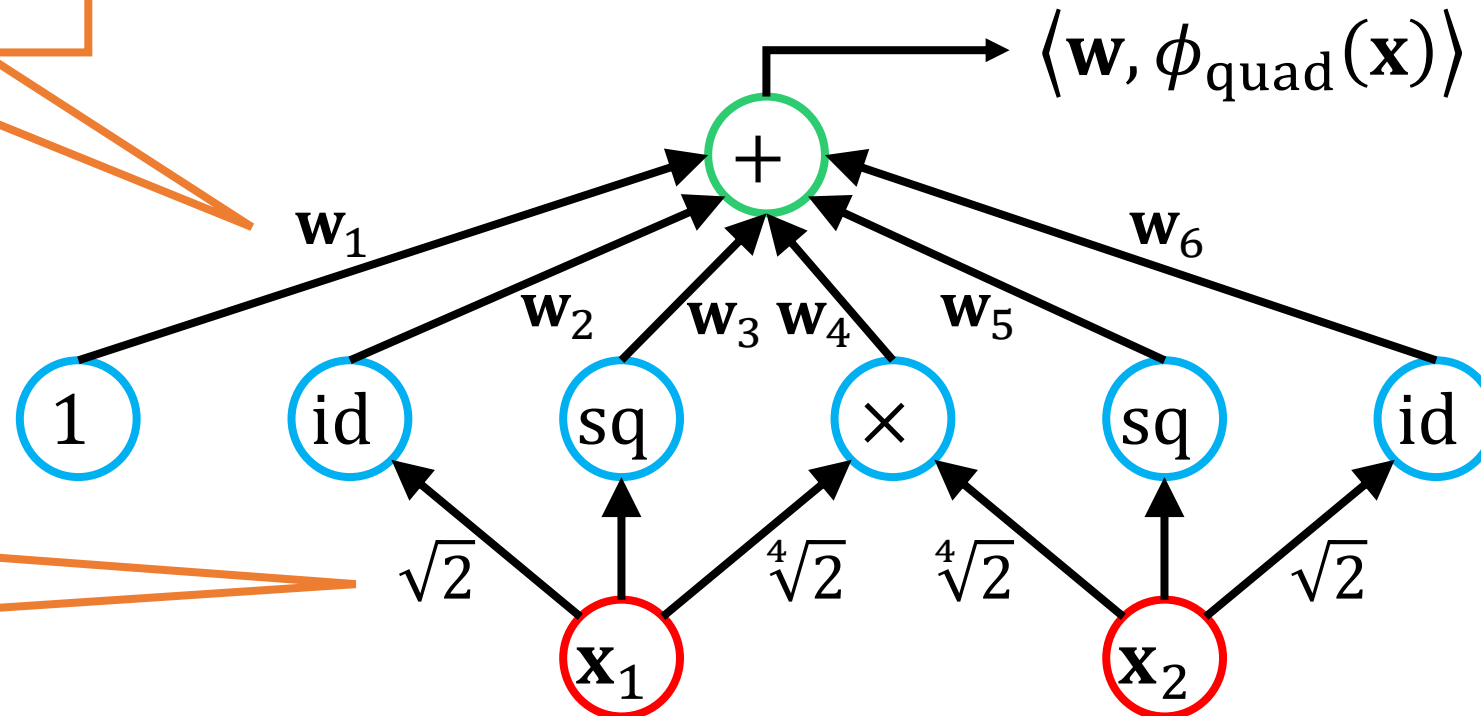
# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The feature map for  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$

$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$

Training an SVM using GD/CD tunes these weights over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$

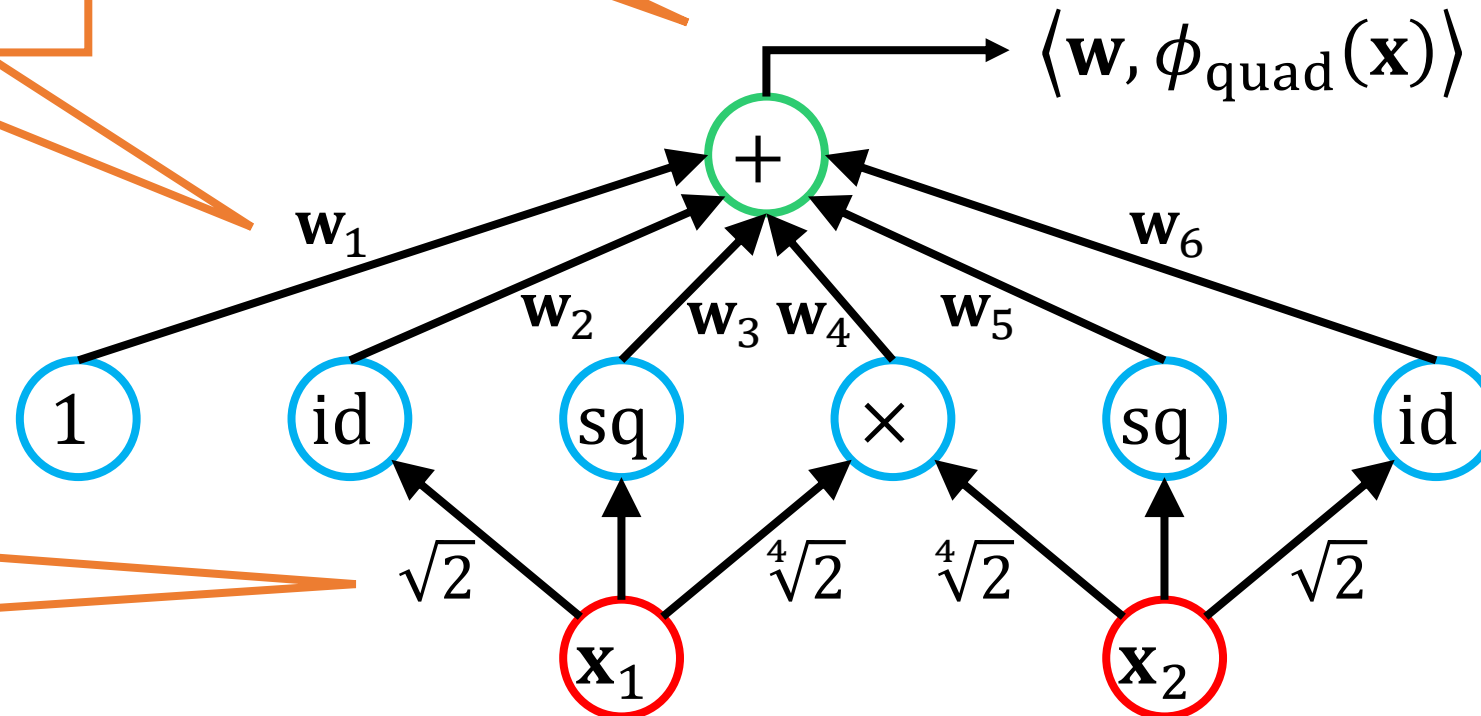


But not these weights

# Back to Kernels first

Can represent any quadratic fn over  $\mathbf{x}$

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- A 3 layer network  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
 $\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$
- Training an SVM using GD/CD tunes these weights over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$



But not these weights

# Back to Kernels first

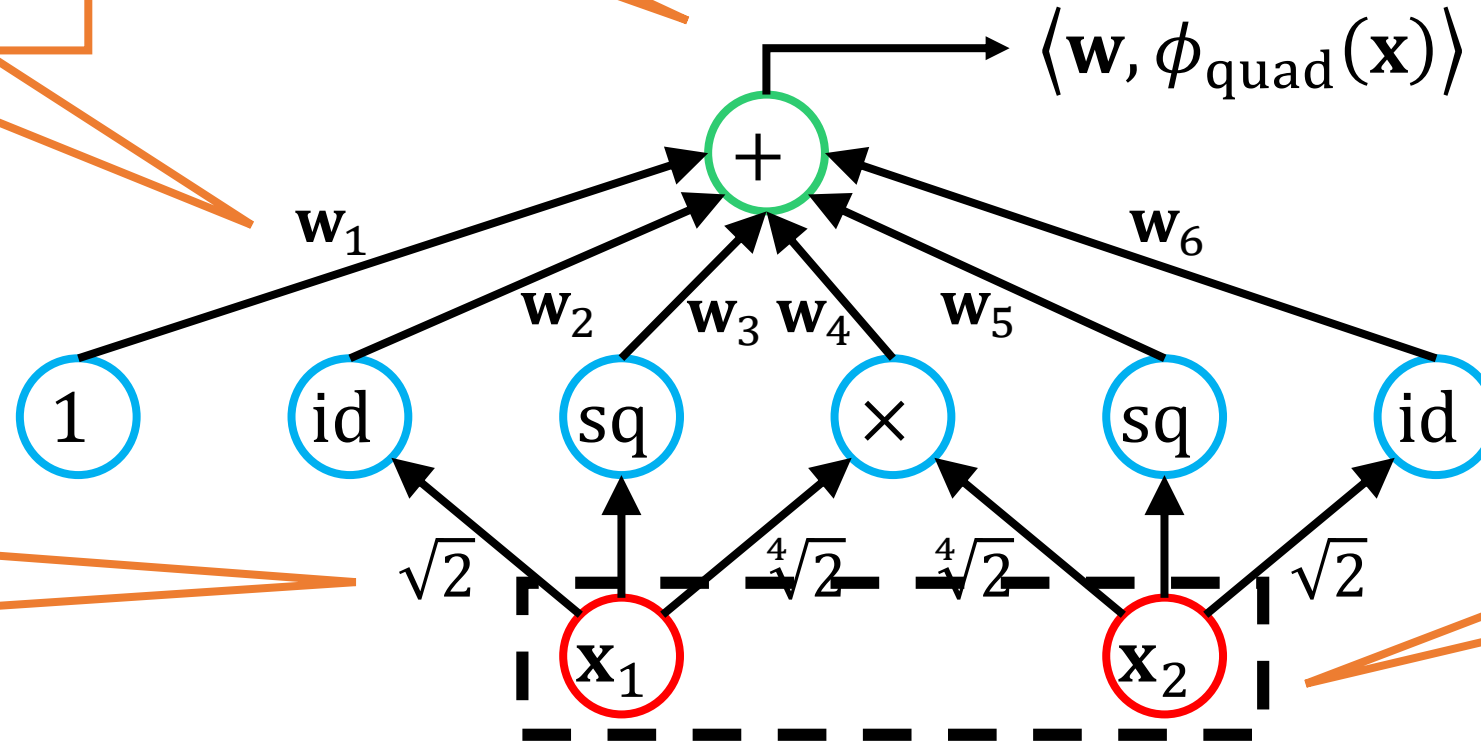
Can represent any quadratic fn over  $\mathbf{x}$

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$

- The A 3 layer network  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$

$$\phi_{\text{quad}}(\mathbf{x}) = [1, \sqrt{2} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$$

Training an SVM using GD/CD tunes over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$  these weights



But not these weights

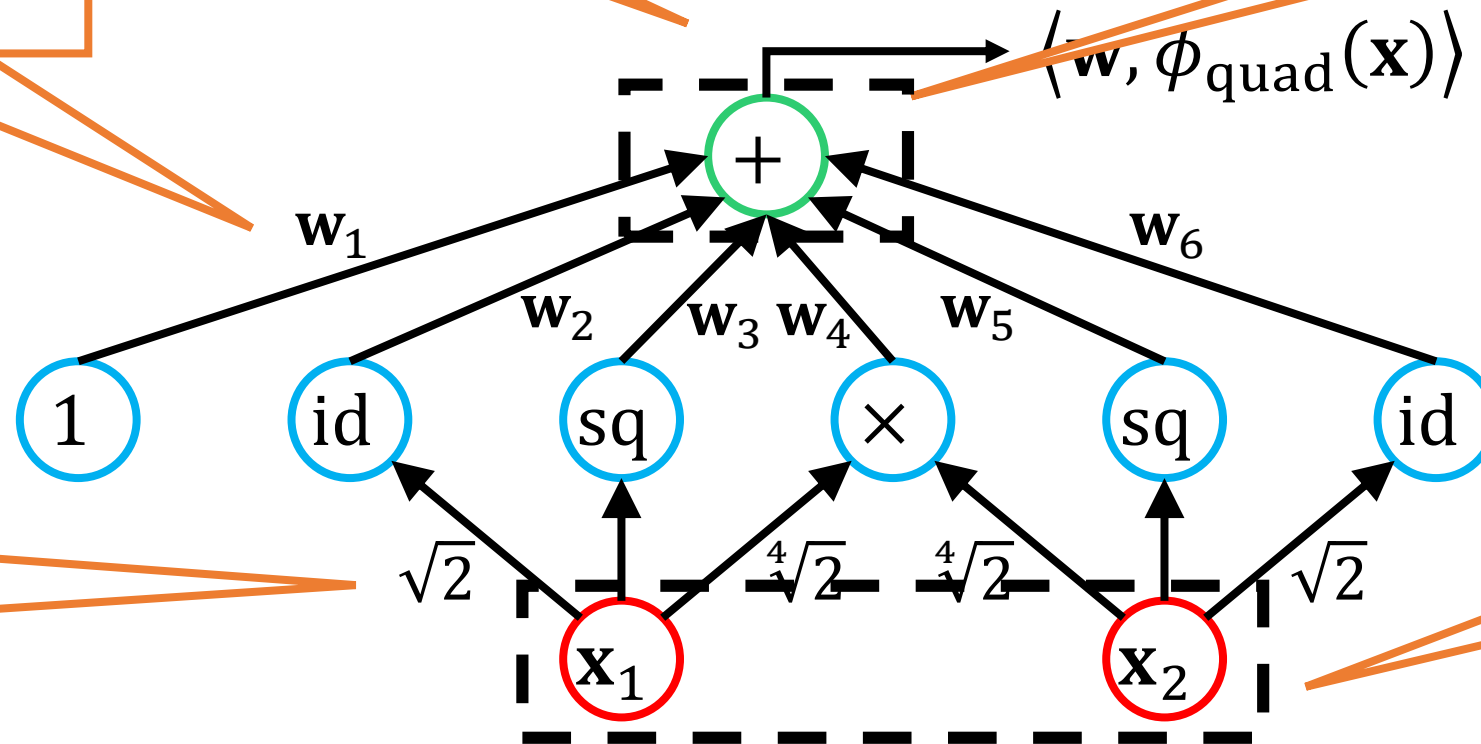
Input layer

# Back to Kernels first

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$

- The A 3 layer network  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$

Training an SVM using GD/CD tunes these weights over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$



Can represent any quadratic fn over  $\mathbf{x}$

Output layer

But not these weights

Input layer

# Back to Kernels first

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$

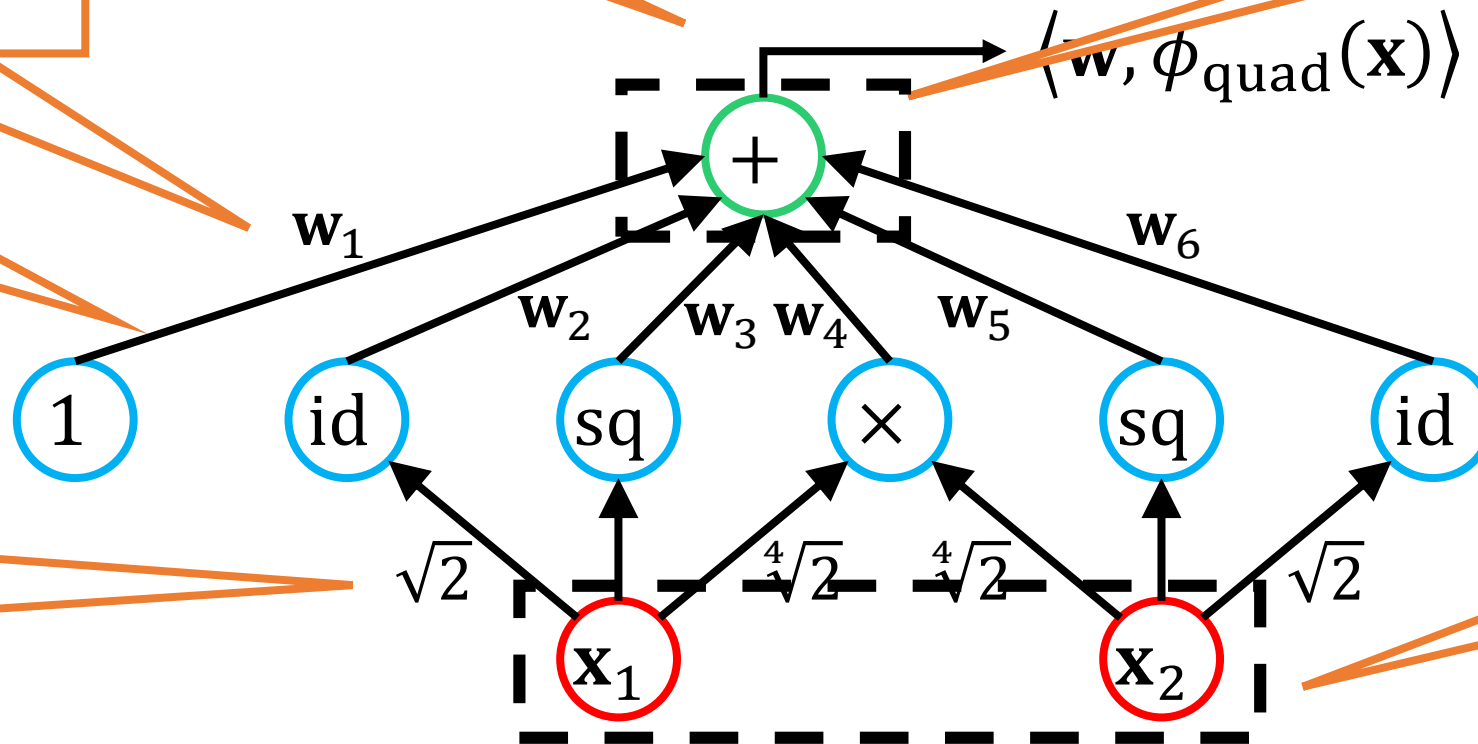
- The A 3 layer network  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$

Training an SVM using GD/CD tunes these weights over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$

Output layer

I/O layers are called "visible"

But not these weights



# Back to Kernels first

- Consider the quadratic kernel  $K_{\text{quad}} = (\langle \mathbf{x}^1, \mathbf{x}^2 \rangle + 1)^2$  on  $\mathcal{X} = \mathbb{R}^2$
- The A 3 layer network  $K_{\text{quad}}$  is  $\phi_{\text{quad}}$  where for  $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2) \in \mathbb{R}^2$   
 $\phi_{\text{quad}}(\mathbf{x}) = [\frac{1}{\sqrt{2}} \cdot \mathbf{x}_1, \mathbf{x}_1^2, \sqrt{2} \cdot \mathbf{x}_1 \cdot \mathbf{x}_2, \mathbf{x}_2^2, \sqrt{2} \cdot \mathbf{x}_2] \in \mathbb{R}^6$
- Training an SVM using GD/CD tunes these weights over  $\phi_{\text{quad}}$  is represented as a vector  $\mathbf{w} \in \mathbb{R}^6$

Can represent any quadratic fn over  $\mathbf{x}$

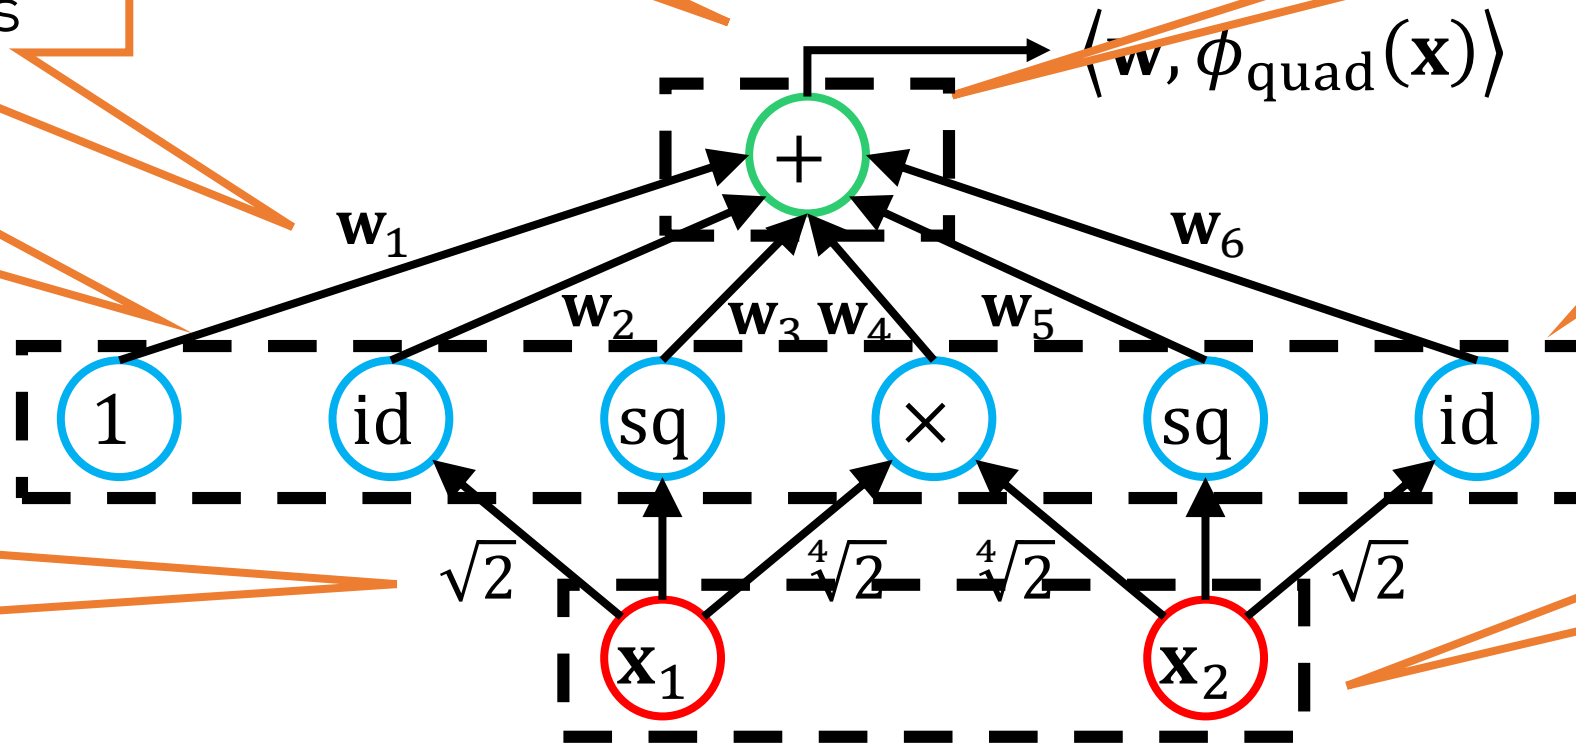
Output layer

Hidden layer

Input layer


I/O layers are called "visible"

But not these weights




# The “neuron” in Neural Networks

# The “neuron” in Neural Networks





# The “neuron” in Neural Networks




$\mathbf{x}_1$

$\mathbf{x}_2$


...

$\mathbf{x}_d$

# The “neuron” in Neural Networks




# The “neuron” in Neural Networks



# The "neuron" in Neural Networks

## Input

Input layer: no i/p  
Hidden/output layer: i/p  
from other neurons




# The "neuron" in Neural Networks

## Input

Input layer: no i/p  
Hidden/output layer: i/p  
from other neurons

Some input items can  
be a constant e.g. 1



# The "neuron" in Neural Networks


## Activation

Input/output layer: id  
Hidden layer: non-linear

## Input

Input layer: no i/p  
Hidden/output layer: i/p  
from other neurons

Some input items can  
be a constant e.g. 1



# in Neural Networks

## Output

Output layer: final o/p  
Input/hidden layer: o/p  
to other neurons


## Activation

Input/output layer: id  
Hidden layer: non-linear

## Input

Input layer: no i/p  
Hidden/output layer: i/p  
from other neurons

Some input items can  
be a constant e.g. 1



# Neural Networks

## Output

Output layer: final o/p  
Input/hidden layer: o/p  
to other neurons

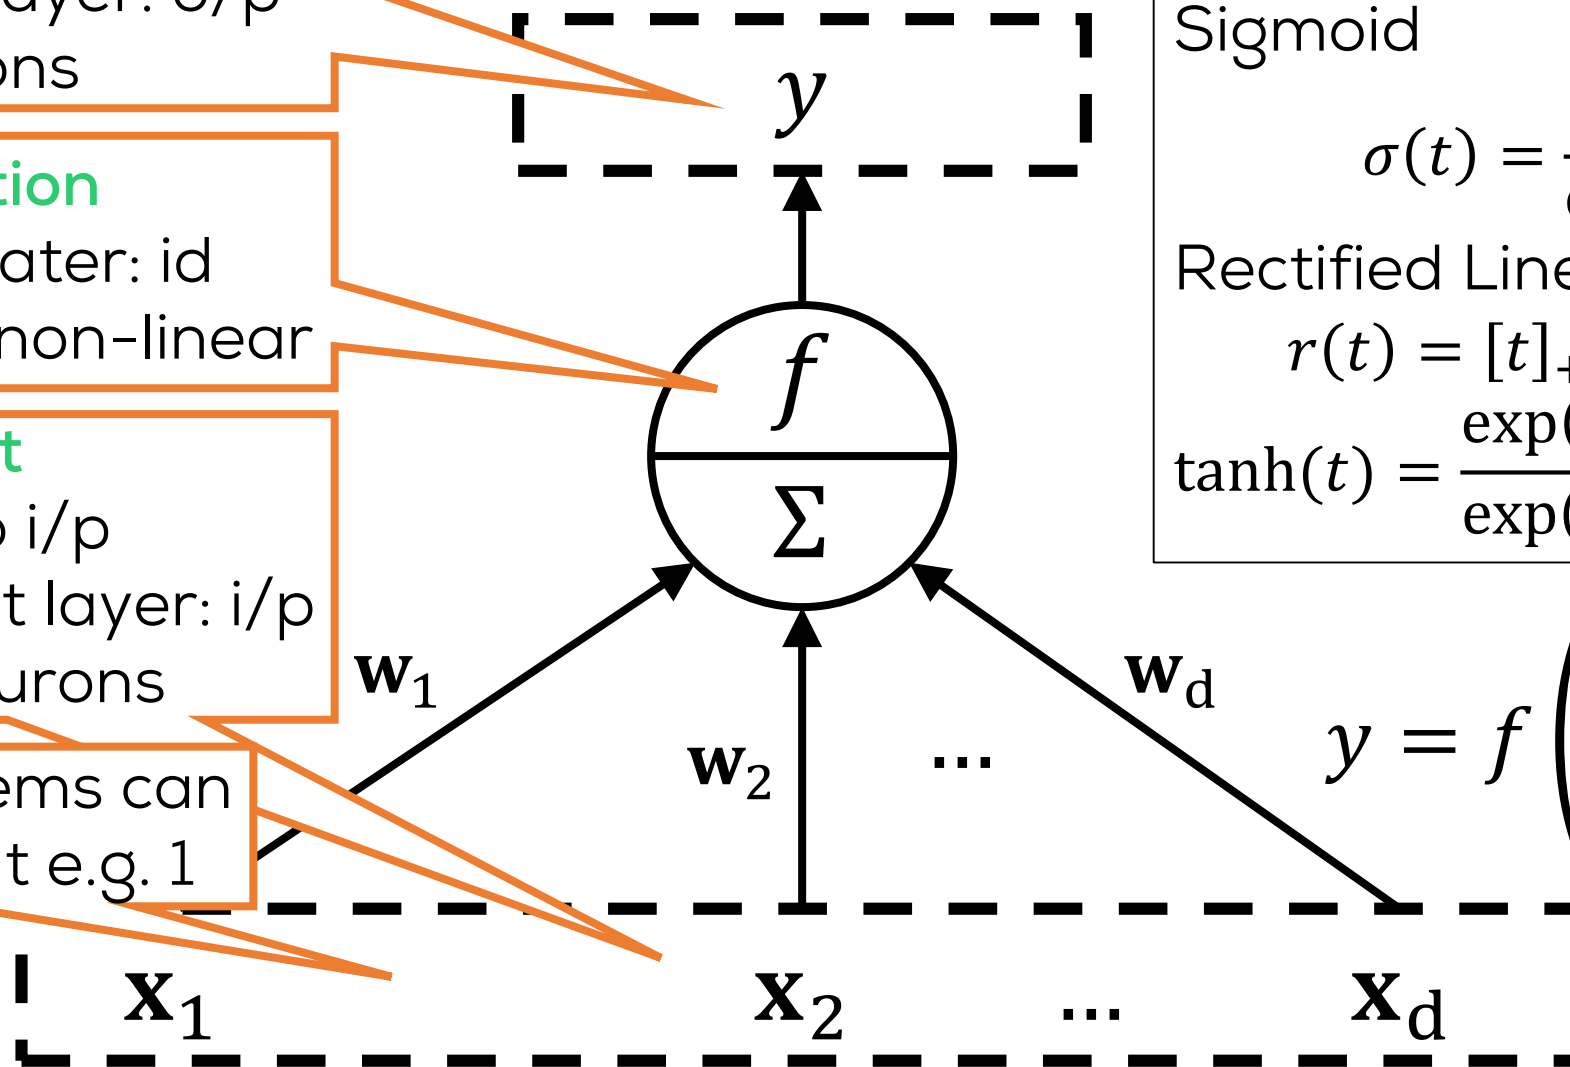
## Activation

Input/output layer: id  
Hidden layer: non-linear

## Input

Input layer: no i/p  
Hidden/output layer: i/p  
from other neurons

Some input items can  
be a constant e.g. 1



## Common "activation" fns $f$

Sigmoid

$$\sigma(t) = \frac{\exp(t)}{\exp(t) + 1}$$

Rectified Linear Unit (ReLU)

$$r(t) = [t]_+ = \max(t, 0)$$

$$\tanh(t) = \frac{\exp(2t) - 1}{\exp(2t) + 1}$$



# Neural Networks

## Output

Output layer: final o/p  
Input/hidden layer: o/p to other neurons

## Activation

Input/output layer: id  
Hidden layer: non-linear

## Input

Input layer: no i/p  
Hidden/output layer: i/p from other neurons

Some input items can be a constant e.g. 1


## Common "activation" fns $f$

Sigmoid

If there is no hidden layer then network is just a linear model, also called a **perceptron**

$$r(t) = [t]_+ = \max(t, 0)$$

$$\tanh(t) = \frac{\exp(2t) - 1}{\exp(2t) + 1}$$



## Output

Output layer: final o/p  
Input/hidden layer: o/p to other neurons

## Activation

Input/output layer: id  
Hidden layer: non-linear

## Input

Input layer: no i/p  
Hidden/output layer: i/p from other neurons


Some input items can be a constant e.g. 1

Sometimes output layer is given a non-id activation. Matter of convention


Common "activation" fns  $f$   
Sigmoid

If there is no hidden layer then network is just a linear model, also called a **perceptron**

$$r(t) = [t]_+ = \max(t, 0)$$
$$\tanh(t) = \frac{\exp(2t) - 1}{\exp(2t) + 1}$$




# A Feedforward Network



# A Feedforward Network


Can I connections  
jump layers?



# A Feedforward Network

Can I connections  
jump layers?


Not in a  
feedforward NN



# A Feedforward Network

Can I connections  
jump layers?

Not in a  
feedforward NN




No "reverse"  
links either

# A Feedforward Network


Can I connections  
jump layers?

Not in a  
feedforward NN



No "reverse"  
links either

# A Feedforward Network



Can I connections jump layers?


Not in a feedforward NN

Can the activation function vary across layers?

No "reverse" links either



# A Feedforward Network



Can I connections  
jump layers?


Not in a  
feedforward NN

Can the activation  
function vary  
across layers?

Yes. Wait  
for CNNs

No "reverse"  
links either

# A Feedforward Network



Can I connections jump layers?

Not in a feedforward NN

One output node needed for binary classfn, regresn, more for multi-label/class

Can the activation function vary across layers?

Yes. Wait for CNNs

No "reverse" links either

# A Feedforward Network

One output node needed for binary classfn, regresn, more for multi-label/class

Can I connections jump layers?


Not in a feedforward NN

Can the activation function vary across layers?

Yes. Wait for CNNs

All weights are learnt

No "reverse" links either



# A Feedforward Network

One output node needed for binary classfn, regresn, more for multi-label/class

Can I connections jump layers?

Not in a feedforward NN


Connections b/w layers is usually "dense" – all pairs are connected

All weights are learnt

Can the activation function vary across layers?

Yes. Wait for CNNs

No "reverse" links either



# A Feedforward Network

Multi-layered  
**perceptron**

Can I connections  
jump layers?

Not in a  
feedforward NN

Connections  
b/w layers is  
usually "dense"  
– all pairs are  
connected


All weights  
are learnt

One output node needed  
for binary classfn, regresn,  
more for multi-label/class

Can the activation  
function vary  
across layers?

Yes. Wait  
for CNNs

No "reverse"  
links either




# A few thoughts

- Just as in kernel slide, lower layers can be interpreted as working very hard to compute useful and informative features
- Last layer exploits all this hard work to learn a good model
$$y = \sum_{i=1}^{k_l} \mathbf{w}_i \cdot \phi(\mathbf{x}_i), k_l = \text{\#nodes in layer previous to output layer}$$
- Note: output is linear in the features computed by lower layers

# A few thoughts

- Just as in l
- very hard t
- Last layer
- $y = \sum_{i=1}^{k_l} \mathbf{w}_i$
- Note: outp




is working

del  
it layer  
~ layers

# A few thoughts

- Just as in k...
- Last layer...
- Note: output...



is working

del

at layer

r layers



# A few thoughts

- Just as in kernel slide, lower layers can be interpreted as working very hard to compute useful and informative features
- Last layer exploits all this hard work to learn a good model
$$y = \sum_{i=1}^{k_l} \mathbf{w}_i \cdot \phi(\mathbf{x}_i), k_l = \text{\#nodes in layer previous to output layer}$$
- Note: output is linear in the features computed by lower layers
- Can have any no. of layers, any no. of nodes in each layer
- Having a linear activation function is useless since the entire network will then just learn a linear function in input
- ReLU networks always learn piecewise linear functions
- Try proving the above two results by induction on number of hidden layers (base case – no hidden layer) as an exercise

# A few thoughts

- Kernel models work with a vast (often infinite) set of features. NN methods try to learn a small set of features from data itself
- Features are non-linear in kernels as well as NNs
- Why can't I have the nice operations of product, squaring, identity as "activation functions" as in the kernel slide?
- A variant called *Sum-product Networks* (SPN) does exactly this
- Neural networks are also *universal*
- A neural network with a single hidden layer with infinitely many nodes or else infinitely many layers each with finitely many nodes can learn any function of the input (details technical)
- Next class: how NNs are trained

# Please give your Feedback

<http://tinyurl.com/ml17-18afb>